

# Artificial Intelligence and Machine Learning

## Unit II

### Linear Regression

#### My own latex definitions

```
In [11]: import matplotlib
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
plt.style.use('seaborn-whitegrid')

font = {'family' : 'Times',
        'weight' : 'bold',
        'size'   : 12}

matplotlib.rc('font', **font)

# Aux functions

def plot_grid(Xs, Ys, axs=None):
    ''' Aux function to plot a grid'''
    t = np.arange(Xs.size) # define progression of int for indexing colormap
    if axs:
        axs.plot(0, 0, marker='*', color='r', linestyle='none') #plot origin
        axs.scatter(Xs,Ys, c=t, cmap='jet', marker='.') # scatter x vs y
        axs.axis('scaled') # axis scaled
    else:
        plt.plot(0, 0, marker='*', color='r', linestyle='none') #plot origin
        plt.scatter(Xs,Ys, c=t, cmap='jet', marker='.') # scatter x vs y
        plt.axis('scaled') # axis scaled

def linear_map(A, Xs, Ys):
    '''Map src points with A'''
    # [NxN,NxN] -> NxNx2 # add 3-rd axis, like adding another layer
    src = np.stack((Xs,Ys), axis=Xs.ndim)
    # flatten first two dimension
    # (NN)x2
    src_r = src.reshape(-1,src.shape[-1]) #ask reshape to keep last dimension and adjust the rest
    # 2x2 @ 2x(NN)
    dst = A @ src_r.T # 2xNN
    #(NN)x2 and then reshape as NxNx2
    dst = (dst.T).reshape(src.shape)
    # Access X and Y
    return dst[...,0], dst[...,1]

def plot_points(ax, Xs, Ys, col='red', unit=None, linestyle='solid'):
    '''Plots points'''
    ax.set_aspect('equal')
    ax.grid(True, which='both')
    ax.axhline(y=0, color='gray', linestyle="--")
    ax.axvline(x=0, color='gray', linestyle="--")
    ax.plot(Xs, Ys, color=col)
    if unit is None:
        plotVectors(ax, [[0,1],[1,0]], ['gray']*2, alpha=1, linestyle=linestyle)
    else:
        plotVectors(ax, unit, [col]*2, alpha=1, linestyle=linestyle)

def plotVectors(ax, vecs, cols, alpha=1, linestyle='solid'):
    '''Plot set of vectors.'''
    for i in range(len(vecs)):
        x = np.concatenate([[0,0], vecs[i]])
        ax.quiver([x[0]],
                  [x[1]],
                  [x[2]],
                  [x[3]],
                  angles='xy', scale_units='xy', scale=1, color=cols[i],
                  alpha=alpha, linestyle=linestyle, linewidth=2)
```

```

In [2]: import matplotlib
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
plt.style.use('seaborn-whitegrid')

font = {'family' : 'Times',
        'weight' : 'bold',
        'size'   : 12}

matplotlib.rc('font', **font)

# Aux functions

def plot_grid(Xs, Ys, axs=None):
    ''' Aux function to plot a grid'''
    t = np.arange(Xs.size) # define progression of int for indexing colormap
    if axs:
        axs.plot(0, 0, marker='*', color='r', linestyle='none') #plot origin
        axs.scatter(Xs,Ys, c=t, cmap='jet', marker='.') # scatter x vs y
        axs.axis('scaled') # axis scaled
    else:
        plt.plot(0, 0, marker='*', color='r', linestyle='none') #plot origin
        plt.scatter(Xs,Ys, c=t, cmap='jet', marker='.') # scatter x vs y
        plt.axis('scaled') # axis scaled

def linear_map(A, Xs, Ys):
    '''Map src points with A'''
    # [NxN,NxN] -> NxNx2 # add 3-rd axis, like adding another layer
    src = np.stack((Xs,Ys), axis=Xs.ndim)
    # flatten first two dimension
    # (NN)x2
    src_r = src.reshape(-1,src.shape[-1]) #ask reshape to keep last dimension and adjust the rest
    # 2x2 @ 2x(NN)
    dst = A @ src_r.T # 2xNN
    # (NN)x2 and then reshape as NxNx2
    dst = (dst.T).reshape(src.shape)
    # Access X and Y
    return dst[...0], dst[...1]

def plot_points(ax, Xs, Ys, col='red', unit=None, linestyle='solid'):
    '''Plots points'''
    ax.set_aspect('equal')
    ax.grid(True, which='both')
    ax.axhline(y=0, color='gray', linestyle="--")
    ax.axvline(x=0, color='gray', linestyle="--")
    ax.plot(Xs, Ys, color=col)
    if unit is None:
        plotVectors(ax, [[0,1],[1,0]], ['gray']*2, alpha=1, linestyle=linestyle)
    else:
        plotVectors(ax, unit, [col]*2, alpha=1, linestyle=linestyle)

def plotVectors(ax, vecs, cols, alpha=1, linestyle='solid'):
    '''Plot set of vectors.'''
    for i in range(len(vecs)):
        x = np.concatenate([[0,0], vecs[i]])
        ax.quiver([x[0]],
                  [x[1]],
                  [x[2]],
                  [x[3]],
                  angles='xy', scale_units='xy', scale=1, color=cols[i],
                  alpha=alpha, linestyle=linestyle, linewidth=2)

```

## Recap previous lecture

- Model Selection and Assessment
- Cross-validation
- Evaluation Metrics

# Today's lecture

We go back to your loved 🧡 Linear Algebra

## Supervised, Parametric Models

- 1) Ordinary Linear Regression with Least Squares
- 2) Probabilistic Interpretation
- 3) Gradient Descent "Family"

## This lecture material is taken from


- [Mostly from Stanford class](#)
- [Stanford notes](#)
- [Tibshirani - Chapter 4 page 43](#)
- [Sklearn model selection](#)
- [Bishop - Chapter 3 page 137](#)

```
In [3]: import numpy as np
from matplotlib import pyplot as plt

from sklearn import linear_model, datasets

n_samples = 100
size = 10

X, y, coef_gt = datasets.make_regression(
    n_samples=n_samples,
    n_features=2,
    n_informative=1,
    noise=20,
    coef=True,
    random_state=42,
)
fig = plt.figure(figsize=(size, size))
ax = fig.add_subplot(projection='3d')
ax.scatter(X[:, 0], X[:, 1], y, c='red', marker='o')
ax.view_init(0, -90)
```




```
In [4]: import numpy as np
from matplotlib import pyplot as plt

from sklearn import linear_model, datasets

n_samples = 100
size = 10

X, y, coef_gt = datasets.make_regression(
    n_samples=n_samples,
    n_features=2,
    n_informative=1,
    noise=20,
    coef=True,
    random_state=42,
)
fig = plt.figure(figsize=(size, size))
ax = fig.add_subplot(projection='3d')
ax.scatter(X[:, 0], X[:, 1], y, c='red', marker='o')
ax.view_init(0, -90)
```



```
In [5]: table = "| x_1 | x_2 | y | \n | --- | --- | --- \n"
for count, (ex, ey) in enumerate(zip(X,y)):
    table += f"| {str(ex[0])[:6]} | {str(ex[1])[:6]} | {str(ey)[:6]} \n"
    if count == 10: break
```

## The data

```
{{print(table)}}
```

## Living area vs Apartment Price

Living area (feet <sup>2</sup> )	Price (1000\$)
2104	400
1600	330
2400	369
1416	232
3000	540
⋮	⋮

# Linear Regression settings

We want to regress  $y$  from  $\mathbf{x}$  that is we want to learn a function  $f_{\theta}$  parametrized by some parameters  $\theta$  so that  $f_{\theta}: \underbrace{\mathbf{x}}_{\text{input}} \mapsto \underbrace{y}_{\text{output}}$

- $\mathbf{x} \in \mathbb{R}^d$  (here  $d=2$ )
- $y$  is a scalar that is continuous  $y \in \mathbb{R}$
- We have a finite number of samples  $D = \{\mathbf{x}_i, y_i\}_{i=1}^n \sim p(\mathbf{x}, y)$  that which labels are generated from a function  $f$  plus error so  $y = f(\mathbf{x}) + \epsilon$

## Linear Hypothesis

We assume relations  $f \longleftrightarrow y$  is **linear**.

We know  $D = \{\mathbf{x}_i, y_i\}_{i=1}^n$  and we want to find  $\theta \doteq (\theta_0, \dots, \theta_d)$

$$f_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 \cdot x_1 + \theta_2 \cdot x_2 + \dots + \theta_d \cdot x_d$$

So  $\theta \doteq (\theta_0, \dots, \theta_d) \in \mathbb{R}^{d+1}$ .

$$f_{\theta}(\mathbf{x}) = \left( \sum_{i=1}^d \theta_i \cdot x_i \right) + \theta_0$$

## Trick for Notation Compactness

We can augment each feature to have a **bias (intercept term)** set to 1 so that  $\mathbf{x} \doteq [1, \mathbf{x}]$ .

Doing so  $\mathbf{x} \in \mathbb{R}^{d+1}$

$$f_{\theta}(\mathbf{x}) = \theta_0 \cdot \underbrace{x_0}_{\text{always 1}} + \theta_1 \cdot x_1 + \theta_2 \cdot x_2 + \dots + \theta_d \cdot x_d$$

So  $\theta \doteq (\theta_0, \dots, \theta_d) \in \mathbb{R}^{d+1}$ .

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^d \theta_i \cdot x_i = \theta^T \mathbf{x}$$

## Parametric Nature

No matter how many training points  $N$  you have, the parameters are fixed in  $\theta$ .

Note that  $\theta \in \mathbb{R}^{d+1}$ .

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^d \theta_i \cdot x_i = \theta^T \mathbf{x}$$

## Loss or Cost Function for Linear Regression

You see now that the loss is more explicit compared to non-parametric models (K-NN, Decision Trees).

$$J(\theta; \mathbf{x}, y) = \frac{1}{2} \sum_{i=1}^n L(y_i, f_{\theta}(\mathbf{x}_i))$$

where

$$L(y, f_{\theta}(\mathbf{x})) = (f_{\theta}(\mathbf{x}) - y)^2$$

The loss is **the squared error**.

```
{{eps = np.arange(-100,100);plt.plot(eps,eps**2);plt.xlabel('Difference');_=plt.ylabel('Cost')}};
```

# Minimize the Total Loss with a Closed Form Solution

We need to minimize

$$J(\theta; \mathbf{x}, y) = \frac{1}{2} \sum_{i=1}^n L(y_i, f_{\theta}(\mathbf{x}_i))$$

so to find:

$$\theta^* = \arg \min_{\theta} J(\theta; \mathbf{x}, y)$$

## Explicit Cost

$$J(\theta; \mathbf{x}, y) = \frac{1}{2} \sum_{i=1}^n (\underbrace{\theta^T \mathbf{x}_i}_{f_{\theta}} - y_i)^2$$

## Vectorizing the Explicit Cost

$$J(\theta; \mathbf{x}, y) = \frac{1}{2} \sum_{i=1}^n (\underbrace{\theta^T \mathbf{x}_i}_{f_{\theta}} - y_i)^2$$

We define the **design matrix**  $\mathbf{X} \in \mathbb{R}^{n \times (d+1)}$  and **label matrix**  $\mathbf{y} \in \mathbb{R}^{n \times 1}$

$$\mathbf{X} = \begin{bmatrix} - & \mathbf{x}_1^T & - \\ - & \mathbf{x}_2^T & - \\ & \vdots & \\ - & \mathbf{x}_n^T & - \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

The parameters  $\theta \in \mathbb{R}^{d+1}$  are:

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_d \end{bmatrix}$$

## Vectorizing the Explicit Cost

$$\underbrace{\mathbf{X}}_{\mathbb{R}^{n \times (d+1)}} \underbrace{\theta}_{\mathbb{R}^{(d+1) \times 1}} - \underbrace{\mathbf{y}}_{\mathbb{R}^n}$$

$$\underbrace{\mathbf{X}\theta}_{\mathbb{R}^n} - \underbrace{\mathbf{y}}_{\mathbb{R}^n}$$

$$\mathbf{X}\theta - \mathbf{y} = \begin{bmatrix} \mathbf{x}_1^T \theta - y_1 \\ \mathbf{x}_2^T \theta - y_2 \\ \vdots \\ \mathbf{x}_n^T \theta - y_n \end{bmatrix}$$

# Vectorizing the Explicit Cost

$$J(\theta; \mathbf{X}, \mathbf{y}) = \frac{1}{2} (\mathbf{X}\theta - \mathbf{y})^T (\mathbf{X}\theta - \mathbf{y}) = \frac{1}{2} \sum_{i=1}^n \underbrace{(\theta^T \mathbf{x}_i - y_i)^2}_{f_{\theta}}$$

## Solve it

Set the gradient to zero to find **critical points**:

$$\nabla_{\theta} J(\theta; \mathbf{X}, \mathbf{y}) = 0$$

$$\nabla_{\theta} \frac{1}{2} (\mathbf{X}\theta - \mathbf{y})^T (\mathbf{X}\theta - \mathbf{y}) = 0$$

For now forget the "equal to zero":

$$\nabla_{\theta} \frac{1}{2} [(\mathbf{X}\theta)^T (\mathbf{X}\theta) - \underbrace{(\mathbf{X}\theta)^T \mathbf{y}}_{\text{scalar}} - \underbrace{\mathbf{y}^T (\mathbf{X}\theta)}_{\text{scalar}} + \mathbf{y}^T \mathbf{y}]$$

$$\nabla_{\theta} \frac{1}{2} [(\mathbf{X}\theta)^T (\mathbf{X}\theta) - 2\theta^T (\mathbf{X}^T \mathbf{y}) + \mathbf{y}^T \mathbf{y}]$$

$$\nabla_{\theta} \frac{1}{2} [\theta^T (\mathbf{X}^T \mathbf{X}) \theta - 2\theta^T (\mathbf{X}^T \mathbf{y}) + \mathbf{y}^T \mathbf{y}]$$

$$\frac{1}{2} [2\mathbf{X}^T \mathbf{X} \theta - 2\mathbf{X}^T \mathbf{y}]$$

## Set the gradient to zero

$$\frac{1}{2} [2\mathbf{X}^T \mathbf{X} \theta - 2\mathbf{X}^T \mathbf{y}] = 0$$

To get the normal equation

$$\mathbf{X}^T \mathbf{X} \theta = \mathbf{X}^T \mathbf{y}$$

## Final Least Squares solution

Assumes  $\mathbf{X}^T \mathbf{X}$  is **invertible**:

$$\theta = (\underbrace{\mathbf{X}^T \mathbf{X}}_{\text{pseudo inverse}})^{-1} \mathbf{X}^T \mathbf{y} = \mathbf{X}^+ \mathbf{y}$$

where:

$$\mathbf{X}^+ \doteq (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$$



```

In [6]: %matplotlib notebook
from sklearn import linear_model, datasets
from matplotlib import pyplot as plt
import numpy as np

n_samples = 100
size = 8

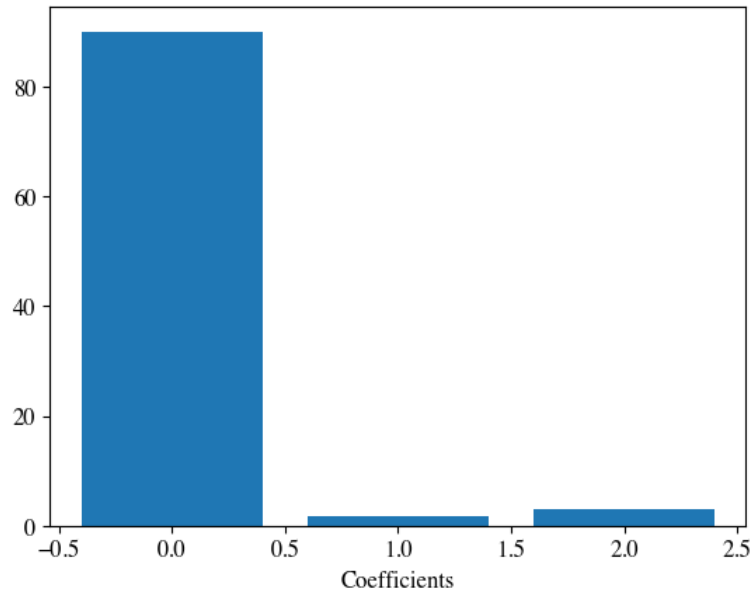
X, y, coef_gt = datasets.make_regression(
    n_samples=n_samples,
    n_features=2,
    n_informative=1,
    noise=20,
    coef=True,
    random_state=42,
)
fig = plt.figure(figsize=(size, size))
ax = fig.add_subplot(projection='3d')
# Linear Regression
bias = np.ones((X.shape[0], 1))
X = np.hstack((X, bias))
theta = np.linalg.inv(X.T@X)@X.T@y
# Now MeshGrid
Xmin, Xmax = X.min(), X.max()
support = np.linspace(Xmin, Xmax, 10)
xx, yy = np.meshgrid(support, support)
data = np.stack((xx, yy), axis=2)
data = data.reshape(-1, 2)
data = np.hstack((data, np.ones((data.shape[0], 1))))
z = np.dot(theta, data.T)
z = z.reshape(xx.shape)
ax.plot_surface(xx, yy, z, alpha=0.2)
ax.scatter(X[:, 0], X[:, 1], y, c='red', marker='o')
ax.view_init(0, 90)

```


## Debugging the Coefficients

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^d \theta_i \cdot x_i = \theta^T \mathbf{x}$$


```
In [7]: %matplotlib inline
plt.figure()
plt.bar(list(range(theta.size)),theta); plt.xlabel('Coefficients');
```




```
In [8]: %matplotlib inline
plt.figure()
plt.bar(list(range(coef_gt.size)),coef_gt); plt.xlabel('Coefficients GT');
```



Important: The distance is NOT orthogonal



Important: The distance is NOT orthogonal (1D case)



Interpretation as solving a overdetermined Linear System ( $n \gg d$ )

Assumes  $\mathbf{X}^T \mathbf{X}$  is **invertible** (full rank) and  $n \gg d$ :

$$\theta = (\underbrace{\mathbf{X}^T \mathbf{X}}_{d \times d})^{-1} \underbrace{\mathbf{X}^T}_{d \times n} \underbrace{\mathbf{y}}_{n \times 1}$$

The normal equation gives you a way to invert  $\mathbf{X}^T \mathbf{X}$

$$\mathbf{X}^T \mathbf{X} \theta = \mathbf{X}^T \mathbf{y}$$

it solves the linear system so that the plane best fit all the points with a trade-off given by the square of the residuals (**least squares**).

What happens if  $n = d + 1$ ?

$$\underbrace{\mathbf{X}}_{n \times n} \theta = \mathbf{y}$$

We can invert it "directly"

$$\theta = X^{-1}y$$

Why?

$$(AB)^{-1} = (B^{-1}A^{-1})$$

then:

$$\theta = (X^T X)^{-1} X^T y = X^{-1} (X^T)^{-1} X^T y = X^{-1} y$$

```
In [91]: %matplotlib notebook
from sklearn import linear_model, datasets
from matplotlib import pyplot as plt
import numpy as np

n_samples = 3
size = 12

X, y, coef_gt = datasets.make_regression(
    n_samples=n_samples,
    n_features=2,
    n_informative=1,
    noise=20,
    coef=True,
    random_state=42,
)
fig = plt.figure(figsize=(size, size))
ax = fig.add_subplot(projection='3d')
ax.scatter(X[:, 0], X[:, 1], y, c='blue', marker='o')
# Linear Regression
bias = np.ones((X.shape[0], 1))
X = np.hstack((X, bias))
theta = np.linalg.inv(X.T @ X) @ X.T @ y
# Now MeshGrid
Xmin, Xmax = X.min(), X.max()
support = np.linspace(Xmin, Xmax, 10)
xx, yy = np.meshgrid(support, support)
data = np.stack((xx, yy), axis=2)
data = data.reshape(-1, 2)
data = np.hstack((data, np.ones((data.shape[0], 1))))
z = np.dot(theta, data.T)
z = z.reshape(xx.shape)
ax.plot_surface(xx, yy, z, alpha=0.2)
ax.scatter(X[:, 0], X[:, 1], y, c='red', marker='o')
ax.view_init(0, 90)
```

## What happens if $n = d$ ?

We see the plane passes exactly through the "training points".

## Probabilistic Interpretation

### Probabilistic Interpretation for Linear Regression

To go probabilistic, we have to make an assumption that each  $y$  is generated linearly but with **additive Gaussian Noise**.

So

$$y_i = \theta^T \mathbf{x}_i + \epsilon$$

where

$$\epsilon \sim \mathcal{N}(0, \sigma^2)$$

- We observe  $(\mathbf{x}_i, y_i)$  but we do not know  $\theta$  and the noise  $\epsilon$ .
- The noise changes from sample to sample but we know it is distributed as Gaussian.

# Probabilistic Interpretation for Linear Regression

To go probabilistic, we have to make an assumption that each  $y$  is generated linearly but with **additive Gaussian Noise**.

So

$$\epsilon = y_i - \theta^T \mathbf{x}_i \sim \mathcal{N}(0, \sigma^2)$$

- We observe  $(\mathbf{x}, y_i)$  but we do not know  $\theta$  and the noise  $\epsilon$ .
- The noise changes from sample to sample but we know it is distributed as Gaussian.

## What does the noise look like?


```
In [101]: %matplotlib notebook
from sklearn import linear_model, datasets
from matplotlib import pyplot as plt
import numpy as np

n_samples = 100
size = 5

X, y, coef_gt = datasets.make_regression(
    n_samples=n_samples,
    n_features=2,
    n_informative=1,
    noise=0,
    coef=True,
    random_state=42,
)

fig = plt.figure(figsize=(size, size))
ax = fig.add_subplot(projection='3d')
# Linear Regression
bias = np.ones((X.shape[0], 1))
X = np.hstack((X, bias))
theta = np.linalg.inv(X.T@X)@X.T@y
# Now MeshGrid
Xmin, Xmax = X.min(), X.max()
support = np.linspace(Xmin, Xmax, 10)
xx, yy = np.meshgrid(support, support)
data = np.stack((xx, yy), axis=2)
data = data.reshape(-1, 2)
data = np.hstack((data, np.ones((data.shape[0], 1))))
z = np.dot(theta, data.T)
z = z.reshape(xx.shape)
ax.plot_surface(xx, yy, z, alpha=0.2)
ax.scatter(X[:, 0], X[:, 1], y, c='red', marker='.')
ax.view_init(0, 90)
ey = y[0]
yn = ey + np.random.randn(20)*10
for yns in yn:
    ax.scatter(X[0, 0], X[0, 1], yns, c='blue', marker='o')
ax.scatter(X[0, 0], X[0, 1], ey, c='green', marker='o');
```

What does the noise look like?



## Probabilistic Interpretation for Linear Regression

To go probabilistic, we have to make an assumption that each  $y$  is generated linearly but with **additive Gaussian Noise**.

So

$$\epsilon_i = y_i - \theta^T \mathbf{x}_i \sim \mathcal{N}(0, \sigma^2)$$

which means the **errors behave IID from a Normal Distribution**.

$$p(\epsilon_i) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{\epsilon_i^2}{2\sigma^2}\right)$$

# Probabilistic Interpretation for Linear Regression

We look at the conditional probability of  $y$  given  $\mathbf{x}$  aka  $p(y | \mathbf{x}; \theta)$ :

$$p(y_i | x_i; \theta) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \theta^T x_i)^2}{2\sigma^2}\right)$$

now is function of  $y$  yet centered on  $\theta^T x_i$ :

$$y_i | x_i; \theta \sim \mathcal{N}(\theta^T x_i, \sigma^2)$$

## Estimate $\theta$ by Maximum Likelihood (MLE)

For a single training point:

$$p(y_i | x_i; \theta) \doteq L(\theta; \mathbf{x}_i, y_i) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - \theta^T x_i)^2}{2\sigma^2}\right)$$

## Estimate $\theta$ by Maximum Likelihood (MLE)

For multiple training point  $\{\mathbf{x}_i, y_i\}_i$ , given IID assumptions on  $\epsilon$  and thus  $y | x$

$$\begin{aligned} L(\theta) &= \prod_{i=1}^n p(y^{(i)} | x^{(i)}; \theta) \\ &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right) \end{aligned}$$

## Maximizing the Log Likelihood (MLE)

$$\begin{aligned} \ell(\theta) &= \log L(\theta) \\ &= \log \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right) \\ &= \sum_{i=1}^n \log \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right) \\ &= n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{\sigma^2} \cdot \frac{1}{2} \sum_{i=1}^n (y^{(i)} - \theta^T x^{(i)})^2 \end{aligned}$$



# Maximizing the Log Likelihood (MLE) equals Minimizing the Squared Loss

(Under the assumption that the errors will distribution as Gaussians)

$$\arg \max_{\theta} n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{\sigma^2} \cdot \frac{1}{2} \sum_{i=1}^n \left( y^{(i)} - \theta^T x^{(i)} \right)^2 \rightarrow \arg \min_{\theta} \frac{1}{2} \sum_{i=1}^n \left( y^{(i)} - \theta^T x^{(i)} \right)^2$$

To summarize: Under the previous probabilistic assumptions on the data, least-squares regression corresponds to finding the maximum likelihood estimate of  $\theta$ . This is thus one set of assumptions under which least-squares regression can be justified as a very natural method that's just doing maximum likelihood estimation.

(Note however that the probabilistic assumptions are by no means necessary for least-squares to be a perfectly good and rational procedure, and there may—and indeed there are—other natural assumptions that can also be used to justify it.)

## Let's assume we could not find a closed form solution but we know how to program plus a bit of calculus, can we still solve Linear Regression?

### We cannot derive a closed form solution...

We need to minimize

$$J(\theta; \mathbf{x}, y) = \frac{1}{2} \sum_{i=1}^n L(y_i, f_{\theta}(\mathbf{x}_i))$$

so to find:

$$\theta^* = \arg \min_{\theta} J(\theta; \mathbf{x}, y)$$

## Closed Form Solution; Iterative Methods

In general, if you can find a closed form solution, that is **the best you can do**.

So if your problem is as simple as inverting a linear system, please \*\*invert a linear system and use pseudo-inverse if you need to!\*\*

In case you cannot derive, we can use **numerical, iterative methods**


## A very simple yet effective Iterative method is Gradient Descent

- This part that we explain now starts to be propaedeutic for **Deep Learning**.

```
In [11]: %matplotlib inline
from sklearn import linear_model, datasets
from matplotlib import pyplot as plt
import numpy as np

n_samples = 100
size = 7


X, y, coef_gt = datasets.make_regression(
    n_samples=n_samples,
    n_features=1,
    n_informative=1,
    noise=20,
    coef=True,
    random_state=42,
)
fig = plt.figure(figsize=(size, size))
ax = fig.add_subplot()
# Linear Regression
bias = np.ones((X.shape[0], 1))
X = np.hstack((X, bias))
theta = np.linalg.inv(X.T@X)@X.T@y
# Now MeshGrid
x_interp = np.linspace(Xmin, Xmax, 100)
x_interp = x_interp.reshape(-1,1)
x_interp = np.c_[x_interp,np.ones_like(x_interp)]
y_interp = np.dot(theta, x_interp.T)
ax.scatter(x_interp[:,0], y_interp, alpha=0.7, marker='.')
ax.scatter(X[:,0], y, c='red', marker='o')
ax.set_xlabel('$x$');
ax.set_ylabel('$y$');
```



```
In [12]: # do mesh grid on possible theta, evaluate the loss and plot the countf
theta_sampling = 50
theta_space = np.linspace(-200, 200, theta_sampling)
xxt, yyt = np.meshgrid(theta_space, theta_space)
xxt = xxt.flatten()
yyt = yyt.flatten()
all_coeff = np.stack((xxt, yyt), axis=1)
loss = 0.5*1
```

```
In [13]: losses = []
for coeff in all_coeff:
    diff = np.dot(X, coeff.T) - y
    losses.append(0.5*np.dot(diff.T, diff))
losses = np.array(losses)
losses = losses.reshape(theta_sampling, theta_sampling)
xxt = xxt.reshape(theta_sampling, theta_sampling)
yyt = yyt.reshape(theta_sampling, theta_sampling)
```

```
In [14]: plt.rcParams['axes.grid'] = False
plt.contourf(xxt, yyt, losses, levels=50, cmap='jet')
plt.colorbar()
plt.scatter(coef_gt, 0, color='red', marker='o', s=50)
plt.axis('scaled')
plt.xlabel('$\theta_1$')
plt.ylabel('$\theta_2$');
```




```
In [15]: %matplotlib notebook
fig = plt.figure(figsize=(7,7))
ax = fig.add_subplot(111, projection='3d')
ax.plot_surface(xxt, yyt, losses)
ax.set_xlabel('$\theta_1$')
ax.set_ylabel('$\theta_2$')
ax.set_zlabel('loss')
plt.show()
```


```
In [16]: %matplotlib inline
from sklearn import linear_model, datasets
from matplotlib import pyplot as plt
import numpy as np

n_samples = 100
size = 7

X, y, coef_gt = datasets.make_regression(
    n_samples=n_samples,
    n_features=1,
    n_informative=1,
    noise=20,
    coef=True,
    random_state=42,
)
fig = plt.figure(figsize=(size, size))
ax = fig.add_subplot()
# Linear Regression
bias = np.ones((X.shape[0], 1))
X = np.hstack((X, bias))
theta = np.linalg.inv(X.T@X)@X.T@y
# Now MeshGrid
x_interp = np.linspace(Xmin, Xmax, 100)
x_interp = x_interp.reshape(-1,1)
x_interp = np.c_[x_interp, np.ones_like(x_interp)]
y_interp = np.dot(theta, x_interp.T)
ax.scatter(x_interp[:,0], y_interp, alpha=0.7, marker='.')
ax.scatter(X[:,0], y, c='red', marker='o')
ax.set_xlabel('$x$');
ax.set_ylabel('$y$');
```



```
In [171]: fig = plt.figure(figsize=(size-1, size-1))
plt.rcParams['axes.grid'] = False
plt.contourf(xxt, yyt, losses, levels=50, cmap='jet');
plt.colorbar()
plt.scatter(coef_gt,0,color='red',marker='o',s=50)
plt.axis('scaled')
plt.xlabel('$\theta_1$')
plt.ylabel('$\theta_2$');
```



Ready for an awesome demo?

In [18]: *## Implementation of Gradient Descent for Logistic Regression*

```
import time
%matplotlib notebook

def get_diff(X, theta, y):
    return X@theta - y[... , np.newaxis]

def get_loss(diff):
    return 0.5*np.dot(diff.T, diff)

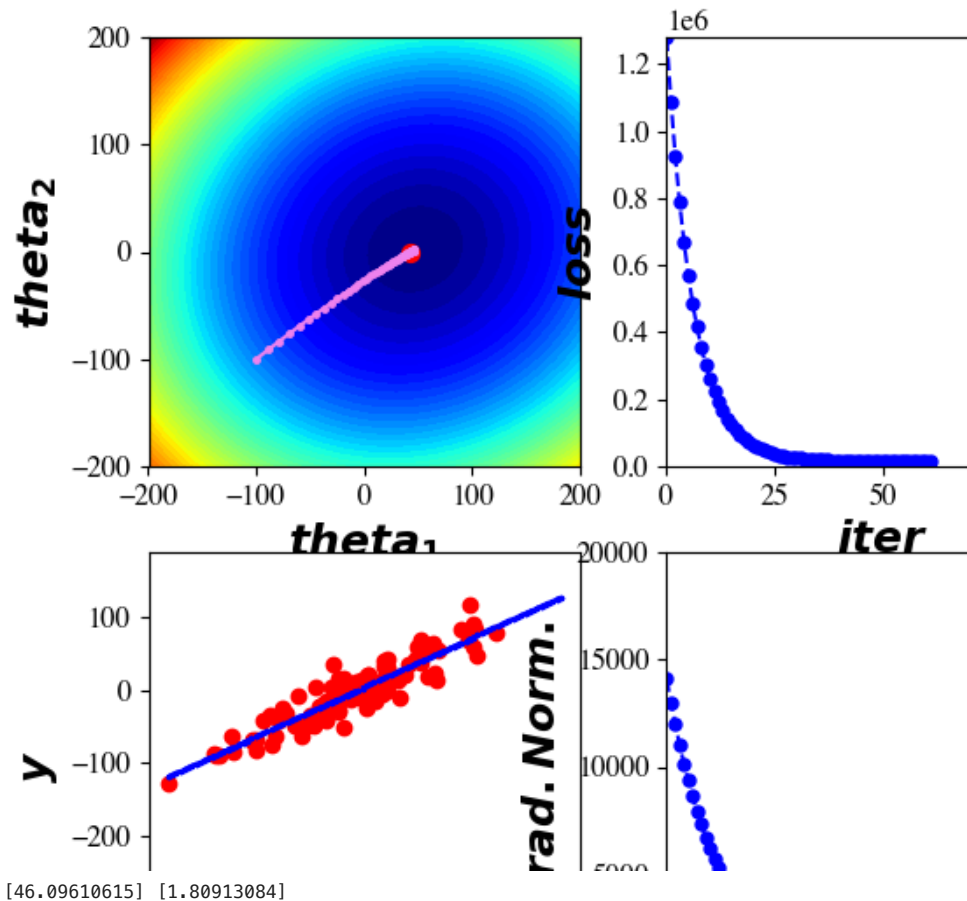
def plot_line(plot3, theta):
    x_interp = np.linspace(Xmin, Xmax, 100)
    x_interp = x_interp.reshape(-1, 1)
    x_interp = np.c_[x_interp, np.ones_like(x_interp)]
    y_interp = np.dot(x_interp, theta)
    if plot3:
        plot3.set_xdata(x_interp[:, 0])
        plot3.set_ydata(y_interp)
    else:
        return x_interp, y_interp

plt.ion()
figure, (axes_1, axes_2) = plt.subplots(2, 2, figsize=(7, 7))
plt.rcParams['axes.grid'] = False
ax0, ax1 = axes_1
ax2, ax3 = axes_2
ax0.contourf(xxt, yyt, losses, levels=50, cmap='jet')
ax0.scatter(coef_gt, 0, color='red', marker='o', s=50)
ax0.set_xlabel('$\theta_1$', fontsize=18)
ax0.set_ylabel('$\theta_2$', fontsize=18)
ax1.set_ylabel('$loss$', fontsize=18)
ax1.set_xlabel('$iter$', fontsize=18)
ax1.set(xlim=(0, 100), ylim=(0, 1.28e6))
ax2.scatter(X[... , 0], y, c='red', marker='o')
ax2.set_xlabel('$x$', fontsize=18)
ax2.set_ylabel('$y$', fontsize=18)
ax3.set_ylabel('$Grad. Norm.$', fontsize=18)
ax3.set_xlabel('$iter$', fontsize=18)
ax3.set(xlim=(0, 100), ylim=(0, 20000))

theta_curr = np.array([[-100, -100]]).T
losses_track = [get_loss(get_diff(X, theta_curr, y))]
grad_norm_track = [1000]

theta_track = np.array(theta_curr)
lr = 1e-3
loss_tol = 10

plot1, = ax0.plot(*theta_curr, color='violet',
                  marker='.', markersize=5, linestyle='--')
plot2, = ax1.plot(*losses_track, color='blue',
                  marker='.', markersize=10, linestyle='--')
xi, yi = plot_line(None, theta_curr)
plot3a, plot3b = ax2.plot(xi, yi, color='blue', marker='.',
                          markersize=3, linestyle='--')
plot4, = ax3.plot(1000, color='blue',
                  marker='.', markersize=10, linestyle='--')
while True:
    diff = get_diff(X, theta_curr, y)
    grad = (diff * X).sum(axis=0, keepdims=True).T
    theta_curr = theta_curr - lr*grad
    theta_track = np.append(theta_track, theta_curr, axis=1)
    diff = get_diff(X, theta_curr, y)
    losses_track.append(get_loss(diff))
    grad_norm_track.append(np.linalg.norm(grad, 2))
    if abs(losses_track[-2]-losses_track[-1]) < loss_tol:
        break
    plot1.set_xdata(theta_track[0, :])
    plot1.set_ydata(theta_track[1, :])
    plot2.set_xdata(range(len(losses_track)))
    plot2.set_ydata(losses_track)
    plot4.set_xdata(range(len(grad_norm_track[1:])))
    plot4.set_ydata(grad_norm_track[1:])
    plot_line(plot3a, theta_curr)
    plot_line(plot3b, theta_curr)
    figure.canvas.draw()
    figure.canvas.flush_events()
    time.sleep(0.1)
print(*theta_curr)
plt.show()
```



## Gradient Descent (GD)

$$J(\theta; \mathbf{x}, y) = \frac{1}{2} \sum_{i=1}^n L(y_i, f_{\theta}(\mathbf{x}_i))$$

where

$$L(y, f_{\theta}(\mathbf{x})) = (f_{\theta}(\mathbf{x}) - y)^2$$


$$\hat{\theta} = \arg \min_{\theta} \frac{1}{2} \sum_{i=1}^n (f_{\theta}(\mathbf{x}_i) - y_i)^2$$

## Analysis

The function  $J(\theta; \mathbf{x}, y)$  is a **convex quadratic function**. The Hessian of  $J(\theta; \mathbf{x}, y)$  at any vector  $\theta$  is the positive definite matrix  $\mathbf{X}^T \mathbf{X}$ . Since  $J$  is lower bounded and grows at infinity, there is a minimum.

- if  $\text{rank}(\mathbf{X}) = \min \{d, n\}$  then  $\mathbf{X}^T \mathbf{X}$  is strictly positive definite. In this case the error function  $J$  is strictly convex, so the **minimum is unique (Ball Shape)**
- if  $\text{rank}(\mathbf{X}) < \min \{d, n\}$  then  $J$  is not strictly convex and the minimum is not unique

```
In [371]: plt.rcParams['axes.grid'] = False
fig = plt.figure(figsize=(4, 4))
ax = fig.add_subplot(111, projection='3d')
ax.plot_surface(xxt, yyt, losses, cmap='jet')
ax.set_xlabel('$\theta_1$')
ax.set_zlabel('$J$')
ax.set_ylabel('$\theta_2$');
```



## Gradient Descent Algorithm as an Iterative Method

**\*\*Idea: make a little step so that locally after each step the cost is lower than before\*\***

Input: Training set  $\{x_i, y_i\}$ , learning rate  $\gamma$ , a small value in  $\{0.1, \dots, 1e-6\}$ .

1. **Initialization - Very Important if the function is not strictly convex**

$$\theta \doteq \mathbf{0}^T$$

Set it to all zeros or random initialization from a distribution. |

1. Repeat until **convergence**:

- Compute the gradient of the loss wrt the parameters  $\theta$  given **all the training set**
- Take a small step in the opposite direction of steepest ascent (**so steepest descent**).

$$\theta \leftarrow \theta - \gamma \nabla_{\theta} J(\theta; \mathbf{x}, y)$$

2. When convergence is reached, your final estimate is in  $\theta$

## Convergence

$$\{\theta_{t=0}, \theta_{t=1}, \dots, \theta_{t=100}\}$$

- 0) Always: validation loss/metric (*early stopping*) (required)

- 1) No significant decrease in the loss function (preferred)

$$|J(\theta; \mathbf{x}, y)_t - J(\theta; \mathbf{x}, y)_{t-1}|$$

- 1) No variations in the parameters

$$||\theta_t - \theta_{t-1}||$$

- 2) Gradient Norm goes to zero

$$||\nabla_{\theta} J(\theta; \mathbf{x}, y)|| \rightarrow 0$$



# Gradient Descent on Linear Regression

1. Initialization

$$\theta \sim \text{random or zero}$$

2.

$$\theta \leftarrow \theta - \gamma \nabla_{\theta} J(\theta; \mathbf{x}, y)$$

$$\theta \leftarrow \theta - \gamma \nabla_{\theta} \frac{1}{2} \sum_{i=1}^n (\theta^T \mathbf{x} - y)^2$$

$$\theta \leftarrow \theta - \gamma \frac{1}{2} \sum_{i=1}^n (2\theta^T \mathbf{x} - 2y) \mathbf{x}$$

$$\theta \leftarrow \theta - \gamma \sum_{i=1}^n (\theta^T \mathbf{x}_i - y_i) \mathbf{x}_i$$

## Dimension Check

Summing up over  $\mathbf{x}_i \in \mathbb{R}^d$  scaled by the difference between the prediction and ground-truth

$$\theta \leftarrow \theta - \gamma \sum_{i=1}^n (\underbrace{\theta^T \mathbf{x}_i}_{\text{scalar}} - \underbrace{y_i}_{\mathbb{R}^d}) \mathbf{x}_i$$

## Stochastic Gradient Descent (SGD)

**Problem of GD:** What happens if  $n \mapsto \infty$ .

To make a small step we have to go through **ALL the training samples**. Optimization could be slow for large  $n$ .

**Idea: make update for each single training sample selected randomly**

$$\theta \leftarrow \theta - \gamma (\underbrace{\theta^T \mathbf{x}_i}_{\text{scalar}} - \underbrace{y_i}_{\mathbb{R}^d}) \mathbf{x}_i \quad \text{where} \quad i \sim \text{U}(0, n)$$

[Many Variations of SGD](#)

Let's see the dynamic of SGD!

Another demo

In [20]: # Implementation of Stochastic Gradient Descent for Logistic Regression

```
import time
%matplotlib notebook

def get_diff(X, theta, y):
    return X@theta - y[:, np.newaxis]

def get_loss(diff):
    return 0.5*np.dot(diff.T, diff)

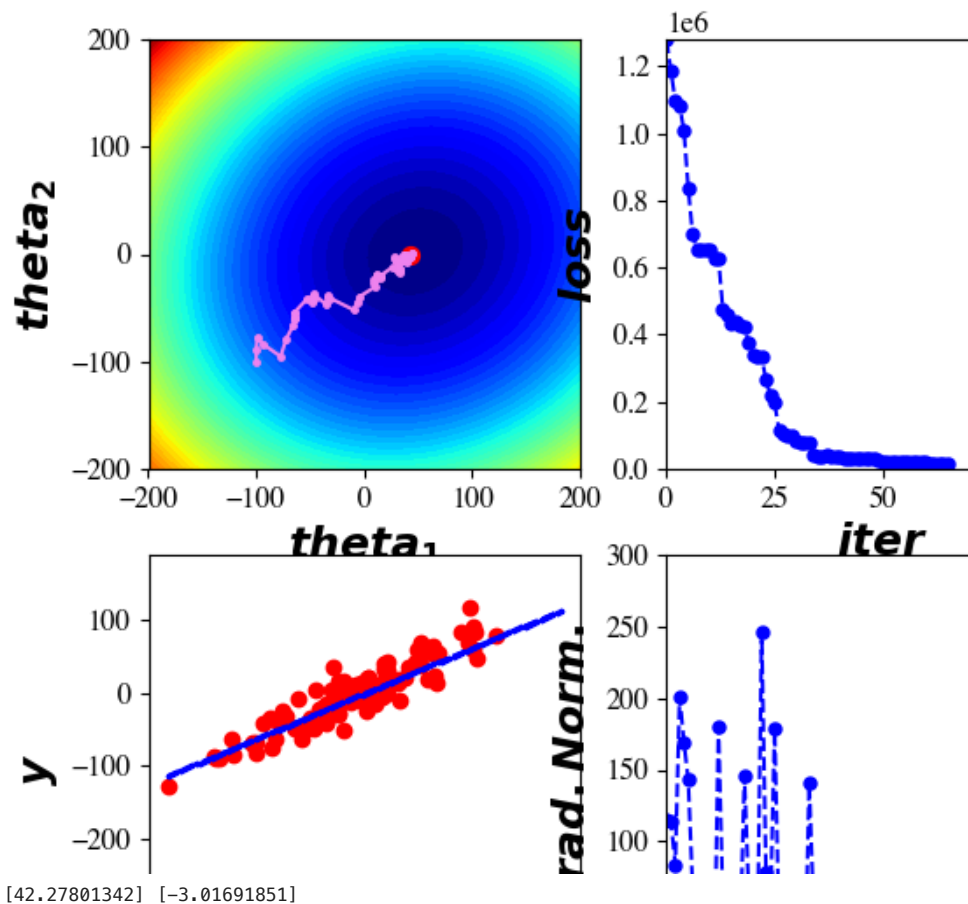
def plot_line(plot3, theta):
    x_interp = np.linspace(Xmin, Xmax, 100)
    x_interp = x_interp.reshape(-1, 1)
    x_interp = np.c_[x_interp, np.ones_like(x_interp)]
    y_interp = np.dot(x_interp, theta)
    if plot3:
        plot3.set_xdata(x_interp[:, 0])
        plot3.set_ydata(y_interp)
    else:
        return x_interp, y_interp

plt.ion()
figure, (axes_1, axes_2) = plt.subplots(2, 2, figsize=(7, 7))
plt.rcParams['axes.grid'] = False
ax0, ax1 = axes_1
ax2, ax3 = axes_2
ax0.contourf(xxt, yyt, losses, levels=50, cmap='jet')
ax0.scatter(coef_gt, 0, color='red', marker='o', s=50)
ax0.set_xlabel('$\theta_1$', fontsize=18)
ax0.set_ylabel('$\theta_2$', fontsize=18)
ax1.set_ylabel('$loss$', fontsize=18)
ax1.set_xlabel('$iter$', fontsize=18)
ax1.set(xlim=(0, 100), ylim=(0, 1.28e6))
ax2.scatter(X[:, 0], y, c='red', marker='o')
ax2.set_xlabel('$x$', fontsize=18)
ax2.set_ylabel('$y$', fontsize=18)
ax3.set_ylabel('$Grad. Norm.$', fontsize=18)
ax3.set_xlabel('$iter$', fontsize=18)
ax3.set(xlim=(0, 100), ylim=(0, 300))

theta_curr = np.array([[-100, -100]]).T
losses_track = [get_loss(get_diff(X, theta_curr, y))]
grad_norm_track = [1000]

theta_track = np.array(theta_curr)
lr = 1e-1
loss_tol = 10
np.random.seed(42)

plot1, = ax0.plot(*theta_curr, color='violet',
                  marker='.', markersize=5, linestyle='--')
plot2, = ax1.plot(*losses_track, color='blue',
                  marker='.', markersize=10, linestyle='--')
xi, yi = plot_line(None, theta_curr)
plot3a, plot3b = ax2.plot(xi, yi, color='blue', marker='.',
                           markersize=3, linestyle='--')
plot4, = ax3.plot(1000, color='blue',
                  marker='.', markersize=10, linestyle='--')
while True:
    diff = get_diff(X, theta_curr, y)
    # STOCHASTIC PART #####
    idx_sampled = np.random.randint(n_samples)
    grad = (diff * X)[idx_sampled, :].T.reshape(-1, 1)
    #####
    theta_curr = theta_curr - lr*grad
    theta_track = np.append(theta_track, theta_curr, axis=1)
    diff = get_diff(X, theta_curr, y)
    losses_track.append(get_loss(diff))
    grad_norm_track.append(np.linalg.norm(grad, 2))
    if abs(losses_track[-2]-losses_track[-1]) < loss_tol:
        break
    plot1.set_xdata(theta_track[0, :])
    plot1.set_ydata(theta_track[1, :])
    plot2.set_xdata(range(len(losses_track)))
    plot2.set_ydata(losses_track)
    plot4.set_xdata(range(len(grad_norm_track[1:])))
    plot4.set_ydata(grad_norm_track[1:])
    plot_line(plot3a, theta_curr)
    plot_line(plot3b, theta_curr)
    figure.canvas.draw()
    figure.canvas.flush_events()
    time.sleep(0.1)
print(*theta_curr)
plt.show()
```



## Stochastic Gradient Descent (SGD): lots of Variations

- Zig-Zag "Noisy" Trajectory of SGD
- Converge to a  $\gamma$ -ball of the solution of GD
- Increases the iterations wrt GD
- Each iteration is so fast that speed of SGD much higher than GD for large training set  $n$ .

## Notable Variations for Deep Learning

- SGD **on mini-batches** (a trade-off between SGD and GD)
- SGD **with momentum** (memory of previous state)

[Many Variations of SGD](#)

## Artificial Intelligence and Machine Learning

### Unit II

### Polynomial Regression, Feature Maps, Ridge Regression

# Recap

We go back to your loved 🧡 Linear Algebra

## Supervised, Parametric Models

- 1) Ordinary Linear Regression with Least Squares
- 2) Probabilistic Interpretation
- 3) Gradient Descent "Family"

## Gradient Descent and [Stochastic] GD

1. Initialization - Very Important if the function is not strictly convex

$$\theta \doteq \mathbf{0}^T$$

Set it to all zeros or random initialization from a distribution.

2. Repeat until **convergence**:

- Compute the gradient of the loss wrt the parameters  $\theta$  given **all the training set**
- Take a small step in the opposite direction of steepest ascent (**so steepest descent**).


$$\theta \leftarrow \theta - \gamma \nabla_{\theta} J(\theta; \mathbf{x}, y)$$

3. When convergence is reached, your final estimate is in  $\theta$

Many Variations of SGD


## GD

$$\theta \leftarrow \theta - \gamma \sum_{i=1}^n (\underbrace{\theta^T \mathbf{x}_i}_{\text{scalar}} - \underbrace{y_i}_{\mathbb{R}^d}) \mathbf{x}_i$$



## SGD

$$\theta \leftarrow \theta - \gamma (\underbrace{\theta^T \mathbf{x}_i}_{\text{scalar}} - \underbrace{y_i}_{\mathbb{R}^d}) \mathbf{x}_i \quad \text{where } i \sim \mathcal{U}(0, n)$$



# Maximizing the Log Likelihood (MLE) equals Minimizing the Squared Loss

(Under the assumption that the errors will distribution as Gaussians)

$$\arg \max_{\theta} n \log \frac{1}{\sqrt{2\pi}\sigma} - \frac{1}{\sigma^2} \cdot \frac{1}{2} \sum_{i=1}^n \left( y^{(i)} - \theta^T x^{(i)} \right)^2 \rightarrow \arg \min_{\theta} \frac{1}{2} \sum_{i=1}^n \left( y^{(i)} - \theta^T x^{(i)} \right)^2$$

To summarize: Under the previous probabilistic assumptions on the data, least-squares regression corresponds to finding the maximum likelihood estimate of  $\theta$ . This is thus one set of assumptions under which least-squares regression can be justified as a very natural method that's just doing maximum likelihood estimation.

(Note however that the probabilistic assumptions are by no means necessary for least-squares to be a perfectly good and rational procedure, and there may—and indeed there are—other natural assumptions that can also be used to justify it.)

## Today

### Make Linear Regression... Non-Linear

### Polynomial Regression with Basis Functions (Feature Map)

### From Feature Maps to Kernel Methods

## This lecture material is taken from

- [Mostly from Bishop - Chapter 3 page 137](#)
- [Stanford notes](#)
- [Tibshirani - Chapter 4 page 43](#)
- [Sklearn Polynomial Regression](#)

## Linear Hypothesis

We assume relations  $f \leftrightarrow y$  is **linear** We know  $D = \{\mathbf{x}_i, y_i\}_{i=1}^n$  and we want to find  $\theta \doteq (\theta_0, \dots, \theta_d)$

$$f_{\theta}(\mathbf{x}) = \theta_0 + \theta_1 \cdot x_1 + \theta_2 \cdot x_2 + \dots + \theta_d \cdot x_d$$

So  $\theta \doteq (\theta_0, \dots, \theta_d) \in \mathbb{R}^{d+1}$ .

$$f_{\theta}(\mathbf{x}) = \left( \sum_{i=1}^d \theta_i \cdot x_i \right) + \theta_0$$

## Trick for Notation Compactness

We can augment each feature to have a **bias (intercept term)** set to 1 so that  $\mathbf{x} \doteq [1, \mathbf{x}]$ .

Doing so  $\mathbf{x} \in \mathbb{R}^{d+1}$

$$f_{\theta}(\mathbf{x}) = \theta_0 \cdot \underbrace{x_0}_{\text{always 1}} + \theta_1 \cdot x_1 + \theta_2 \cdot x_2 + \dots + \theta_d \cdot x_d$$

So  $\theta \doteq (\theta_0, \dots, \theta_d) \in \mathbb{R}^{d+1}$ .

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^d \theta_i \cdot x_i = \theta^T \mathbf{x}$$

## Linear Function of parameters $\theta$ and input $\mathbf{x}$

With  $\mathbf{x} = [1, x_1, \dots, x_d]$  and  $\theta = [\theta_0, \theta_1, \dots, \theta_d]$ , we have:

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^d \theta_i \cdot x_i = \theta^T \mathbf{x}$$

Does not capture non-linear relationships between  $y$  and  $\mathbf{x}$

## Interpreting Linear Regression with Basis Functions

We can have another dimensionality  $m$  instead of  $d$  by using **Basis Functions**  $\phi(\mathbf{x})$ .

With  $\phi(\mathbf{x}) = [1, \phi(x_1), \dots, \phi(x_m)]$  and  $\theta = [\theta_0, \theta_1, \dots, \theta_m]$ , we have:

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^m \theta_i \cdot x_i = \theta^T \phi(\mathbf{x})$$

For Linear Regression:

- $m = d$ , and
- the **Basis Functions** is :  $\phi_m(\mathbf{x}) = x_m$

## What if..

- $m \neq d$ , and
- the **Basis Functions** is :  $\phi_m(\mathbf{x}) \neq x_m$

then we do not have Linear Regression, and the settings impact the model

## Think Basis Function as (Non-Linear) Transform or Feature Map

Let's consider the one dimensional case. We have already seen that we could change the feature  $x$  to add the bias term  $\mathbf{x} \doteq [x, 1]$

$$\mathbf{x} = \begin{bmatrix} x \\ 1 \end{bmatrix} \rightarrow \phi(\mathbf{x}) = \begin{bmatrix} x^2 \\ x \\ 1 \end{bmatrix}$$

So input dimension is  $d = 2$  then output dimension after  $\phi(\cdot)$  is  $m = 3$ .

In this case we used a second order polynomial to lift up the features

$$\phi_m(\mathbf{x}) = x^m$$

## Basis Function as Non-Linear Transform

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^m \theta_i \cdot \phi_i(x) = \theta^T \phi(\mathbf{x}) = \theta_0 + \theta_1 \cdot x + \theta_2 \cdot x^2$$

**Important observation:**

- This is still **linear function** of the parameters  $\theta$ , in fact we still **take the dot product**
- Though it is **NON linear function** of the features  $x$

## Two Observations

- This is still **linear function** of the parameters  $\theta$ 
  - **Good** we can solve it with Linear Regression
- Though it is **NON linear function** of the features  $x$ 
  - **Even better**, we capture non-linearity in the data

## Polynomial Regression (Basis Function $\phi_m(\mathbf{x}) = x^m$ )

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^m \theta_i \cdot \phi_i(x) = \theta^T \phi(\mathbf{x}) = \theta_0 + \theta_1 \cdot x + \theta_2 \cdot x^2$$

### Important observation:

- This is still **linear function** of the parameters  $\theta$ , in fact we still **take the dot product**
- Though it is **NON linear function** of the features  $x$
- $m$  is the degree of the **Polynomial** we consider
- $m$  the higher the degree, the more expressive is the model

## Poly (Multi) Nomial (Names or Terms)

```
In [211]: sigma_noise = 0.5
support_X = 10
offset_valid = 7
np.random.seed(0)

def gen_data(x, sigma):
    n_samples = x.shape[0]
    return x*np.sin(x) + sigma*np.random.randn(n_samples)

XX = np.random.uniform(-support_X, support_X, size=n_samples)
x = XX[:80]
x_valid = np.random.uniform(-support_X-offset_valid, support_X+offset_valid, size=20)#XX[80:]
y = gen_data(x, sigma=sigma_noise).reshape(-1,1)
y_valid = gen_data(x_valid, sigma=sigma_noise).reshape(-1,1)
x = x.reshape(-1,1)
x_valid = x_valid.reshape(-1,1)
```

```
In [221]: plt.figure(figsize=(7, 7))
_ = plt.scatter(x, y, c='red', marker='.')
_ = plt.scatter(x_valid, y_valid, c='blue', marker='.')
plt.legend(['Train', 'Valid']);
```

```
In [231]: %matplotlib inline
```

## Non-Linear Data


### Linear Hypothesis? 🤖

```
{{plt.figure(figsize=(7,7)); plt.scatter(x,y, c='red', marker='o', s=30);_=plt.scatter(x_valid,y_valid, c='blue', marker='o', s=30);}}
```

## Now let's solve it with Linear Regression (we assume no bias)

```
In [241]: def solve_lstq(x, y):  
    return np.linalg.inv(x.T@x)@x.T@y  
  
plt.figure(figsize=(7, 7))  
plt.scatter(x, y, c='red', marker='o', s=30)  
plt.scatter(x_valid, y_valid, c='blue', marker='o', s=30)  
theta = solve_lstq(x, y)  
x_interp = np.linspace(-support_X*1.5, support_X*1.5, 100).reshape(-1, 1)  
y_interp = np.dot(theta, x_interp.T)  
plt.plot(x_interp[:, 0], y_interp.T, alpha=0.7, marker='.');
```





## Let's check the training error (or fitting error)

```
err = np.power(y - np.dot(theta, x.T), 2).mean()
```

Numerically it seems pretty high!

```
{{print(np.power(y - np.dot(theta, x.T), 2).mean())}}
```

```
In [25]: errors = []
errors.append(np.power(y - np.dot(theta, x.T), 2).mean())
```

## Let's try a quadratic basis function

Let's consider the one dimensional case. We have already seen that we could change the feature  $x$  to add the bias term  $\mathbf{x} \doteq [x, 1]$

$$\mathbf{x} = \begin{bmatrix} x \\ 1 \end{bmatrix} \rightarrow \phi(\mathbf{x}) = \begin{bmatrix} x^2 \\ x \\ 1 \end{bmatrix}$$

So input dimension is  $d = 1$  then output dimension after  $\phi(\cdot)$  is  $m = 2$ .

In this case we used a second order polynomial to lift up the features

$$\phi_m(\mathbf{x}) = x^m$$

## The blessing of dimensionality

Let's consider the one dimensional case. We have already seen that we could change the feature  $x$  to add the bias term  $\mathbf{x} \doteq [x, 1]$

$$\mathbf{x} = \begin{bmatrix} x \\ 1 \end{bmatrix} \rightarrow \phi(\mathbf{x}) = \begin{bmatrix} x^2 \\ x \\ 1 \end{bmatrix}$$

- In some sense this can be seen as opposite to the **curse of dimensionality**
- Though if you increase  $m$  too much you may face **curse of dimensionality** again

```
In [26]: %matplotlib notebook
xq = np.c_[x, x**2] # make quadratic features
fig = plt.figure(figsize=(size, size))
ax = fig.add_subplot(projection='3d')
ax.scatter(xq[:, 0], xq[:, 1], y, c='red', marker='o', s=30)
ax.view_init(0, -90)
```

## Now we can still solve it with LS but $m = 2$

We can have another dimensionality  $m$  instead of  $d$  by using **Basis Functions**  $\phi(\mathbf{x})$ .


With  $\phi(\mathbf{x}) = [1, \phi(x_1), \dots, \phi(x_m)]$  and  $\theta = [\theta_0, \theta_1, \dots, \theta_m]$ , we have:

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^m \theta_i \cdot \phi_i(x) = \theta^T \phi(\mathbf{x})$$

For Linear Regression:

- $m > d$ , and
- the **Basis Functions** is :  $\phi_m(\mathbf{x}) = x^m$

```
In [27]: %matplotlib inline
plt.figure(figsize=(7, 7))
plt.scatter(x, y, c='red', marker='o', s=30)
theta_q = solve_lstq(xq, y)
x_interp = np.linspace(-support_X*1.5, support_X*1.5, 100).reshape(-1, 1)
x_interp_q = np.c_[x_interp, x_interp**2]
y_interp_q = np.dot(theta_q.T, x_interp_q.T)
plt.plot(x_interp_q[:, 0], y_interp_q.T, alpha=0.7, marker='.');
```



Let's check the training error (or fitting error) again

```
err = np.power(y - np.dot(theta_q.T, xq.T), 2).mean()
```

Numerically it seems pretty high!

```
{{print(np.power(y - np.dot(theta_q.T, xq.T), 2).mean())}}
```

```
In [28]: errors.append(np.power(y - np.dot(theta_q.T, xq.T), 2).mean())
```

Now we can still solve it with LS but  $m = 3$

We can have another dimensionality  $m$  instead of  $d$  by using **Basis Functions**  $\phi(\mathbf{x})$ .


With  $\phi(\mathbf{x}) = [1, \phi(x_1), \dots, \phi(x_m)]$  and  $\theta = [\theta_0, \theta_1, \dots, \theta_m]$ , we have:

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^m \theta_i \cdot \phi_i(x) = \theta^T \phi(\mathbf{x})$$

For Linear Regression:

- $m > d$ , and
- the **Basis Functions** is :  $\phi_m(\mathbf{x}) = x^m$

```
In [29]: %matplotlib inline
xc = np.c_[x, x**2, x**3] # make cubic features
plt.figure(figsize=(7, 7))
plt.scatter(x, y, c='red', marker='o', s=30);
theta_c = solve_lstq(xc,y)
x_interp_c = np.c_[x_interp, x_interp**2, x_interp**3]
y_interp_c = np.dot(theta_c.T, x_interp_c.T)
plt.plot(x_interp_c[:, 0], y_interp_c.T, alpha=0.7, marker='.');
```



Let's check the training error (or fitting error) again

```
err = np.power(y - np.dot(theta_c.T, xc.T), 2).mean()
```

Numerically it seems pretty high!

```
{{print(np.power(y - np.dot(theta_c.T, xc.T), 2).mean())}}
```

```
In [38]: errors.append(np.power(y - np.dot(theta_c.T, xc.T), 2).mean())
```

We can analyze what happens in function of  $m$

We can have another dimensionality  $m$  instead of  $d$  by using **Basis Functions**  $\phi(\mathbf{x})$ .

With  $\phi(\mathbf{x}) = [1, \phi(x_1), \dots, \phi(x_m)]$  and  $\theta = [\theta_0, \theta_1, \dots, \theta_m]$ , we have:

$$f_{\theta}(\mathbf{x}) = \sum_{i=0}^m \theta_i \cdot x_i = \theta^T \phi(\mathbf{x})$$

For Linear Regression:


- $m > d$ , and
- the **Basis Functions** is :  $\phi_m(\mathbf{x}) = x^m$



```



In [31]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import Pipeline
import numpy as np



errors = []
errors_valid = []
models = []
for m in range(1, 20):
    model = Pipeline([('poly', PolynomialFeatures(degree=m, include_bias=False, interaction_only=False)),
                      ('linear', LinearRegression(fit_intercept=False))])
    models.append(model)
    model = model.fit(x, y)
    x_interp = np.linspace(-support_X+offset_valid,
                          support_X+offset_valid, 100).reshape(-1, 1)
    y_interp = model.predict(x_interp)
    y_est = model.predict(x)
    errors.append(np.power(y - y_est, 2).mean())
    errors_valid.append(np.power(y_valid - model.predict(x_valid), 2).mean())
    # Draw
    plt.figure(figsize=(7, 7))
    plt.scatter(x, y, c='red', marker='o', s=30)
    plt.plot(x_interp, y_interp, alpha=0.7, marker='.')
    plt.ylim([-10, 10])

```






































Now let's check both the errors in train and validation

```
In [321]: fig, axes = plt.subplots(1,2, figsize=(20, 7))
axes[0].bar(range(1, 20), errors)
axes[1].bar(range(1, 20), errors_valid);
axes[1].set_ylim([0,200])
axes[0].set_ylim([0,200])
m_best = np.argmin(errors_valid)
print(f'M best (polynomial degree is) {m_best+1}')
```


M best (polynomial degree is) 3



```
In 1331: # Draw
plt.figure(figsize=(7, 7))
plt.scatter(x, y, c='red', marker='.')
plt.scatter(x_valid, y_valid, c='blue', marker='.')
y_interp = models[m_best].predict(x_interp)
plt.plot(x_interp, y_interp, alpha=0.7, marker='.');
plt.ylim([-100,100]);
```



```
In 1341: # Draw
plt.figure(figsize=(7, 7))
plt.scatter(x, y, c='red', marker='.')
plt.scatter(x_valid, y_valid, c='blue', marker='.')
y_interp = models[-1].predict(x_interp)
plt.plot(x_interp, y_interp, alpha=0.7, marker='.');
plt.ylim([-100,100]);
```



## Over or Under Fitting


	Underfitting	Optimal	Overfitting
Regression			
Classification			




# Problem of Polynomial Regression

The basis function  $\phi_m(x) = x^m$  is **global** wrt the domain of the feature  $x$ .


- The big limitation of polynomial basis functions is that they are **global functions of the input variable**, so that **changes in one region of input space affect all other regions**.
- This can be resolved by dividing the input space up into regions and fit a different polynomial in each region, leading to **spline functions** (that we do not cover).



## Hint on Spline Functions



## Other Basis Functions



**Figure 3.1** Examples of basis functions, showing polynomials on the left, Gaussians of the form (3.4) in the centre, and sigmoidal of the form (3.5) on the right.

## Limitations of Basis Functions

**Advantage:** It is simple, and your problem stays convex and well behaved. (i.e. you can still use your original gradient descent code, just with the higher dimensional representation)


**Disadvantage:**  $\phi(x)$  might be **Very** high dimensional.

```
X = PolynomialFeatures(interaction_only=True).fit_transform(X)
```

```
In [35]: X_rand = np.array([[3, 7]], dtype=float)
print(f'Input Dimension {X_rand.shape}')
X_rand_poly = PolynomialFeatures(
    degree=2, interaction_only=False).fit_transform(X_rand)
print(f'Output Dimension {X_rand_poly.shape}')
print(X_rand[0, :], X_rand_poly[0, :], sep='\n')
```

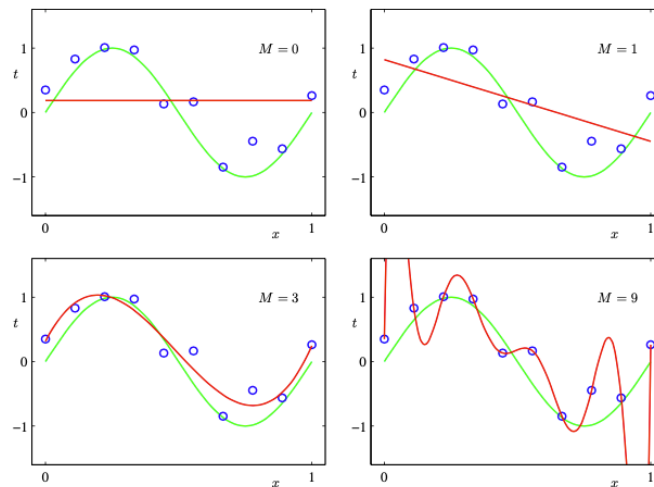
```
Input Dimension (1, 2)
Output Dimension (1, 6)
[3. 7.]
[ 1.  3.  7.  9. 21. 49.]
```

```
In [361]: %matplotlib inline
new_size = [PolynomialFeatures(degree=2, interaction_only=False).fit_transform(
    np.random.rand(1, d)).shape[1] for d in range(1, 50)]
plt.figure(figsize=(12,12))
plt.bar(range(1, 50), new_size);
```



# Debug the Coefficients

## 1.1. Example: Polynomial Curve Fitting 7



**Figure 1.4** Plots of polynomials having various orders  $M$ , shown as red curves, fitted to the data set shown in Figure 1.2.

## Debug the Coefficients → Large Coefficients lead to overfit

**Table 1.1** Table of the coefficients  $w^*$  for polynomials of various order. Observe how the typical magnitude of the coefficients increases dramatically as the order of the polynomial increases.

	$M = 0$	$M = 1$	$M = 6$	$M = 9$
$w_0^*$	0.19	0.82	0.31	0.35
$w_1^*$		-1.27	7.99	232.37
$w_2^*$			-25.43	-5321.83
$w_3^*$			17.37	48568.31
$w_4^*$				-231639.30
$w_5^*$				640042.26
$w_6^*$				-1061800.52
$w_7^*$				1042400.18
$w_8^*$				-557682.99
$w_9^*$				125201.43

## Remedy for Large Coefficients: *Weight Decay* ( $\ell_2$ Regularization)

We minimize the cost plus **penalty term**

$$J(\theta; \mathbf{x}, y) = \frac{1}{2} \sum_{i=1}^n L(y_i, f_{\theta}(\mathbf{x}_i)) + \frac{\lambda}{2} \theta^T \theta = \frac{1}{2} \sum_{i=1}^n L(y_i, f_{\theta}(\mathbf{x}_i)) + \frac{\lambda}{2} \|\theta\|_2^2$$

so to find:

$$\theta^* = \arg \min_{\theta} J_{\text{data}}(\theta; \mathbf{x}, y) + \lambda J_{\text{reg.}}(\theta)$$

## Still has a Closed Form Solution (Regularized Least Squares)

$$\theta = (\lambda \mathbf{I} + \mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Regularization allows complex models to be trained on data sets of limited size without severe over-fitting, essentially by limiting the effective model complexity

## Linear Regression with *Weight Decay* ( $\ell_2$ Regularization)

Consider the eigendecomposition of the symmetric **Positive Semi Definite (PSD)** matrix  $\mathbf{X}\mathbf{X}^T$ :

$$\mathbf{X}^T\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{U}^T = \mathbf{U} \begin{bmatrix} \sigma_1^2 & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \sigma_d^2 \end{bmatrix} \mathbf{U}^T$$

$\underbrace{\hspace{10em}}$   
 $\text{diag} \left( \sigma_1^2, \dots, \sigma_d^2 \right)$

$\sigma_1^2, \dots, \sigma_d^2$  are the eigenvalues and  $\mathbf{U}\mathbf{U}^T = \mathbf{Id}$  (since  $\mathbf{\Sigma}$  is symmetric and square).

If  $\mathbf{X}$  and  $\mathbf{X}\mathbf{X}^T$  are not full rank then some of  $\sigma$  maybe zeros.

## Linear Regression with *Weight Decay* ( $\ell_2$ Regularization)

This implies that if we regularized it, then the pseudo inverse is **always invertible and has a unique solution** since  $\lambda > 0$ :

$$\mathbf{X}^T\mathbf{X} + \lambda \mathbf{I} = \mathbf{U} \begin{bmatrix} \sigma_1^2 + \lambda & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \sigma_d^2 + \lambda \end{bmatrix} \mathbf{U}^T$$

## Linear Regression with *Weight Decay* ( $\ell_2$ Regularization)

This implies that if we regularized it then the pseudo inverse is **always invertible and has a unique solution** since  $\lambda > 0$ :

$$\left( \mathbf{X}^T\mathbf{X} + \lambda \mathbf{I} \right)^{-1} = \mathbf{U} \begin{bmatrix} \frac{1}{\sigma_1^2 + \lambda} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \frac{1}{\sigma_d^2 + \lambda} \end{bmatrix} \mathbf{U}^T$$

## Bias-Variance for Linear Reg. with *Weight Decay* ( $\ell_2$ Regularization)

$$\begin{aligned} \hat{\theta}_n &= \left( \mathbf{X}^T\mathbf{X} + \lambda \mathbf{I} \right)^{-1} \mathbf{X}^T \vec{y} \\ &= \left( \mathbf{X}^T\mathbf{X} + \lambda \mathbf{I} \right)^{-1} \mathbf{X}^T \left( \mathbf{X} \theta^* + \vec{\epsilon} \right) \\ &= \left[ \left( \mathbf{X}^T\mathbf{X} + \lambda \mathbf{I} \right)^{-1} \mathbf{X}^T \mathbf{X} \right] \theta^* + \left[ \left( \mathbf{X}^T\mathbf{X} + \lambda \mathbf{I} \right)^{-1} \mathbf{X}^T \right] \vec{\epsilon} \end{aligned}$$

$$\begin{aligned}
\mathbb{E}[\hat{\theta}_n] &= \mathbb{E}\left[\left[(X^T X + \lambda I)^{-1} X^T X\right] \theta^* + \left[(X^T X + \lambda I)^{-1} X^T\right] \vec{\epsilon}\right] \\
&= \left[(X^T X + \lambda I)^{-1} X^T X\right] \theta^* + \left[(X^T X + \lambda I)^{-1} X^T\right] \mathbb{E}[\vec{\epsilon}] \\
&= \left[(X^T X + \lambda I)^{-1} X^T X\right] \theta^* + \left[(X^T X + \lambda I)^{-1} X^T\right] \vec{0} \\
&= \left[(X^T X + \lambda I)^{-1} X^T X\right] \theta^*
\end{aligned}$$

$$= U \begin{bmatrix} \frac{1}{\sigma_1^2 + \lambda} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \frac{1}{\sigma_d^2 + \lambda} \end{bmatrix} U^T X^T X \theta^*$$

$$= U \begin{bmatrix} \frac{1}{\sigma_1^2 + \lambda} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \frac{1}{\sigma_d^2 + \lambda} \end{bmatrix} \underbrace{U^T U}_{Id} \begin{bmatrix} \sigma_1^2 & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \sigma_d^2 \end{bmatrix} U^T \theta^*$$

$$= U \begin{bmatrix} \frac{1}{\sigma_1^2 + \lambda} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \frac{1}{\sigma_d^2 + \lambda} \end{bmatrix} \begin{bmatrix} \sigma_1^2 & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \sigma_d^2 \end{bmatrix} U^T \theta^*$$

$$= U \begin{bmatrix} \frac{\sigma_1^2}{\sigma_1^2 + \lambda} & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \frac{\sigma_d^2}{\sigma_d^2 + \lambda} \end{bmatrix} U^T \theta^*$$

## Bias-Variance for Linear Reg. with *Weight Decay* ( $\ell_2$ Regularization)

- From the above, we can make a few observations. First, when  $\lambda = 0$ , we see that  $\mathbb{E}[\theta_n] = \theta$ . This implies that **standard linear regression estimator (without regularization) is Unbiased**.
- The more regularization we add (i.e. larger  $\lambda$ ), the smaller the eigenvalues will be, and hence the stronger the "shrinkage" towards 0. Thus it is biased towards zero.


## Weight Decay arises from Bayesian Linear Regression

$$\epsilon = y_i - \theta^T \mathbf{x}_i \sim \mathcal{N}(0, \sigma^2)$$

$$\theta \sim \mathcal{N}(0, \Theta^2) \quad \text{prior on weights}$$

Unlike MLE, now there is a Gaussian prior on the weights; We solve it by doing Maximum A Posteriori (MAP)

## Hint on Bayesian Linear Regression



**Figure 3.7** Illustration of sequential Bayesian learning for a simple linear model of the form  $y(x, \mathbf{w}) = w_0 + w_1 x$ . A detailed description of this figure is given in the text.