


NEURAL COMPUTER GRAPHICS

Antonio D'Orazio





OUTLINE

1. Background: Computer Graphics

1. Applications
2. Terminology
3. Rendering algorithms
4. Ray tracing
5. Representing shapes
6. Implicit representations

2. Neural Computer Graphics

1. Neural implicit representations
2. Training a neural shape
3. Adding details
4. Final pipeline

3. My research



BACKGROUND: COMPUTER GRAPHICS

APPLICATIONS



Videogames

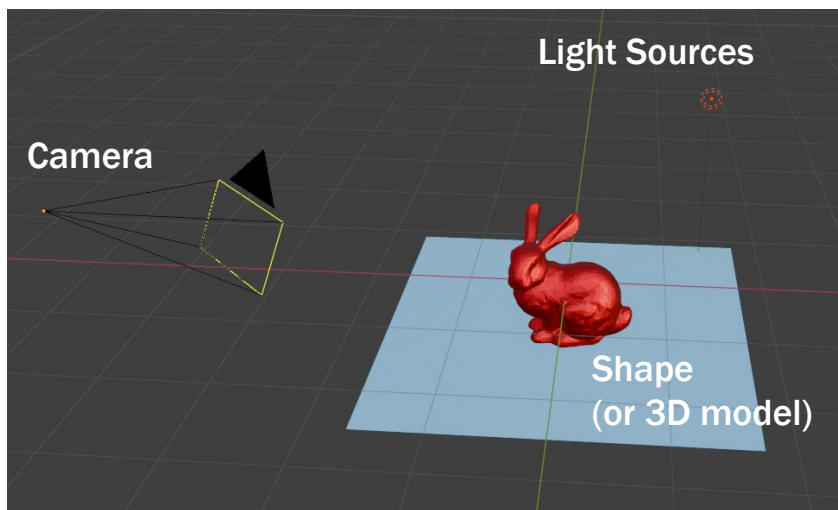


Movies



Architecture


TERMINOLOGY



3D Scene




Rendering



Final 2D Image


RENDERING ALGORITHMS

- Rasterization
 - Fast, approximated
- Ray Tracing
 - Slow, photorealistic




RAY TRACING

- Cast rays from the camera through each pixel




RAY TRACING

- Cast rays from the camera through each pixel
- Find intersections (slow)




RAY TRACING

- Cast rays from the camera through each pixel
- Find intersections (slow)
- Sum the color to the corresponding pixel




RAY TRACING

- Cast rays from the camera through each pixel
- Find intersections (slow)
- Sum the color to the corresponding pixel
- Compute the next bounce(s)




RAY TRACING


- Cast rays from the camera through each pixel
- Find intersections (slow)
- Sum the color to the corresponding pixel
- Compute the next bounce(s)
- Sum shadow and reflection colors to the pixel



REPRESENTING SHAPES



Explicit: mesh



Implicit: signed distance functions (SDFs)

Shape representation: **Explicit** vs **Implicit**

- **Explicit** representations: the 3D model «as is»
 - Graph $G = \{V, E\}$, Set of edges and vertices
 - Hard to manipulate
- **Implicit** representations: we only know how far the model is from us.
 - Signed Distance Functions: $d = f(x, y, z)$
 - When rendering, we use the distances to discover its properties.


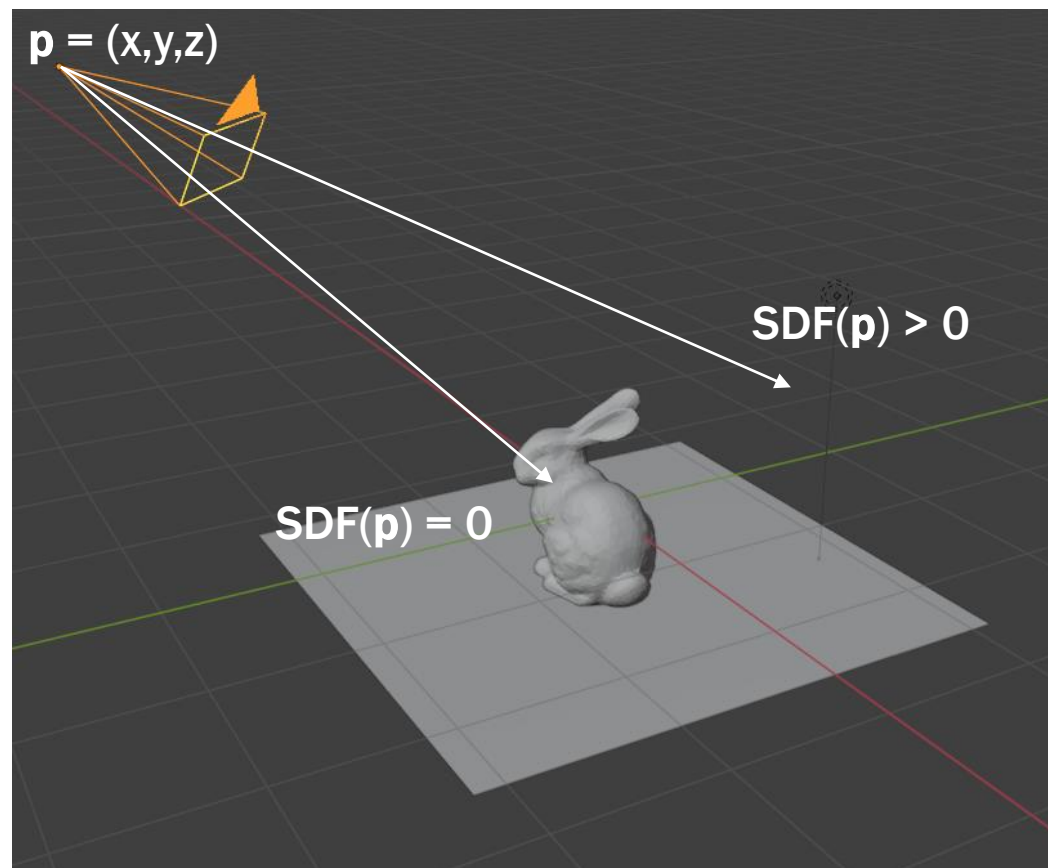
IMPLICIT REPRESENTATIONS

Work well with ray tracing

Q: How do we detect intersections?

A: If $SDF(p) = 0$

We only need to evaluate a function



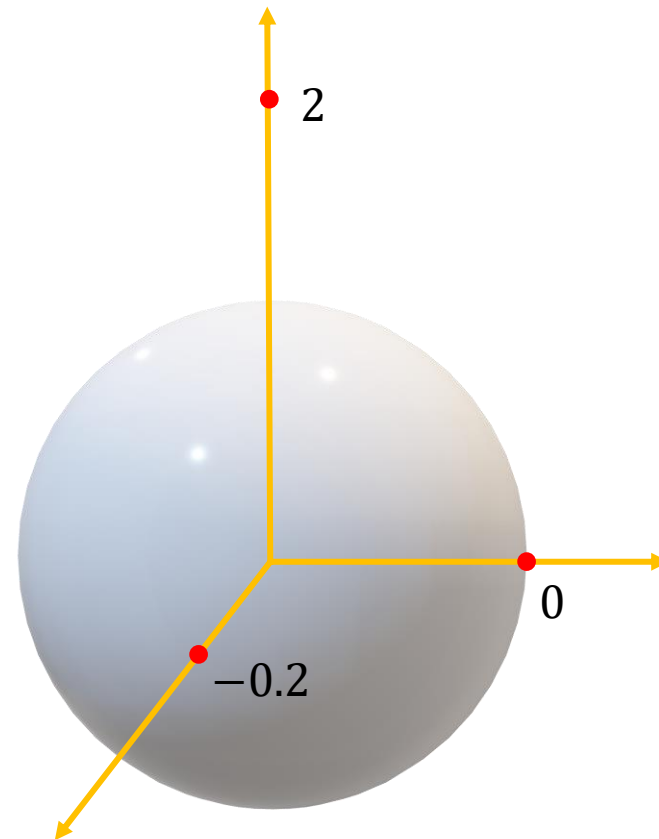
EXAMPLE: IMPLICIT SPHERE

$$SDF(x, y, z) = \sqrt{x^2 + y^2 + z^2} - 1$$

$$SDF(1, 0, 0) = 0 \quad \text{On the surface}$$

$$SDF(0, 0, 0.5) = -0.5 \quad \text{Inside}$$

$$SDF(0, 3, 0) = 2 \quad \text{Outside}$$



A digital rendering of Michelangelo's David statue, featuring a vibrant blue and orange color scheme. The statue is shown from the chest up, with its head turned slightly to the right. The background is a stylized, low-angle view of a classical building with columns and arches, also rendered in the same color palette. The lighting is dramatic, with strong highlights and shadows.


NEURAL GRAPHICS

NEURAL IMPLICIT REPRESENTATIONS

SDFs are functions!

- We can learn them with a neural network

$$SDF_{sphere}(x, y, z) = \sqrt{x^2 + y^2 + z^2} - 1 \quad \approx \quad MLP_{\theta}(x, y, z)$$



- We can learn **very complex** SDFs


NEURAL IMPLICIT REPRESENTATIONS

SDFs are functions!


- We can learn them with a neural network

$$SDF_{sphere}(x, y, z) = \sqrt{x^2 + y^2 + z^2} - 1 \quad \approx \quad MLP_{\theta}(x, y, z)$$


- We can learn very complex SDFs




TRAINING A NEURAL SHAPE



Input mesh



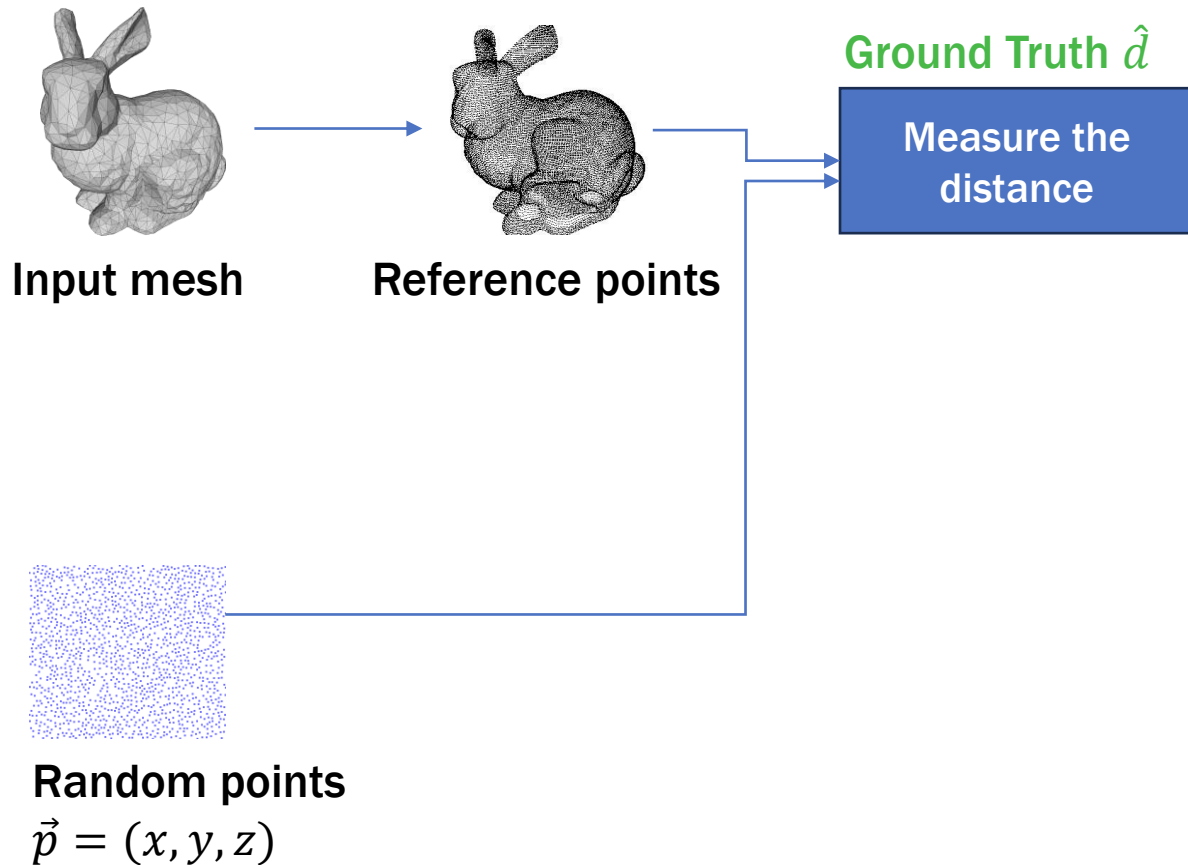
Reference points




Random points

$$\vec{p} = (x, y, z)$$

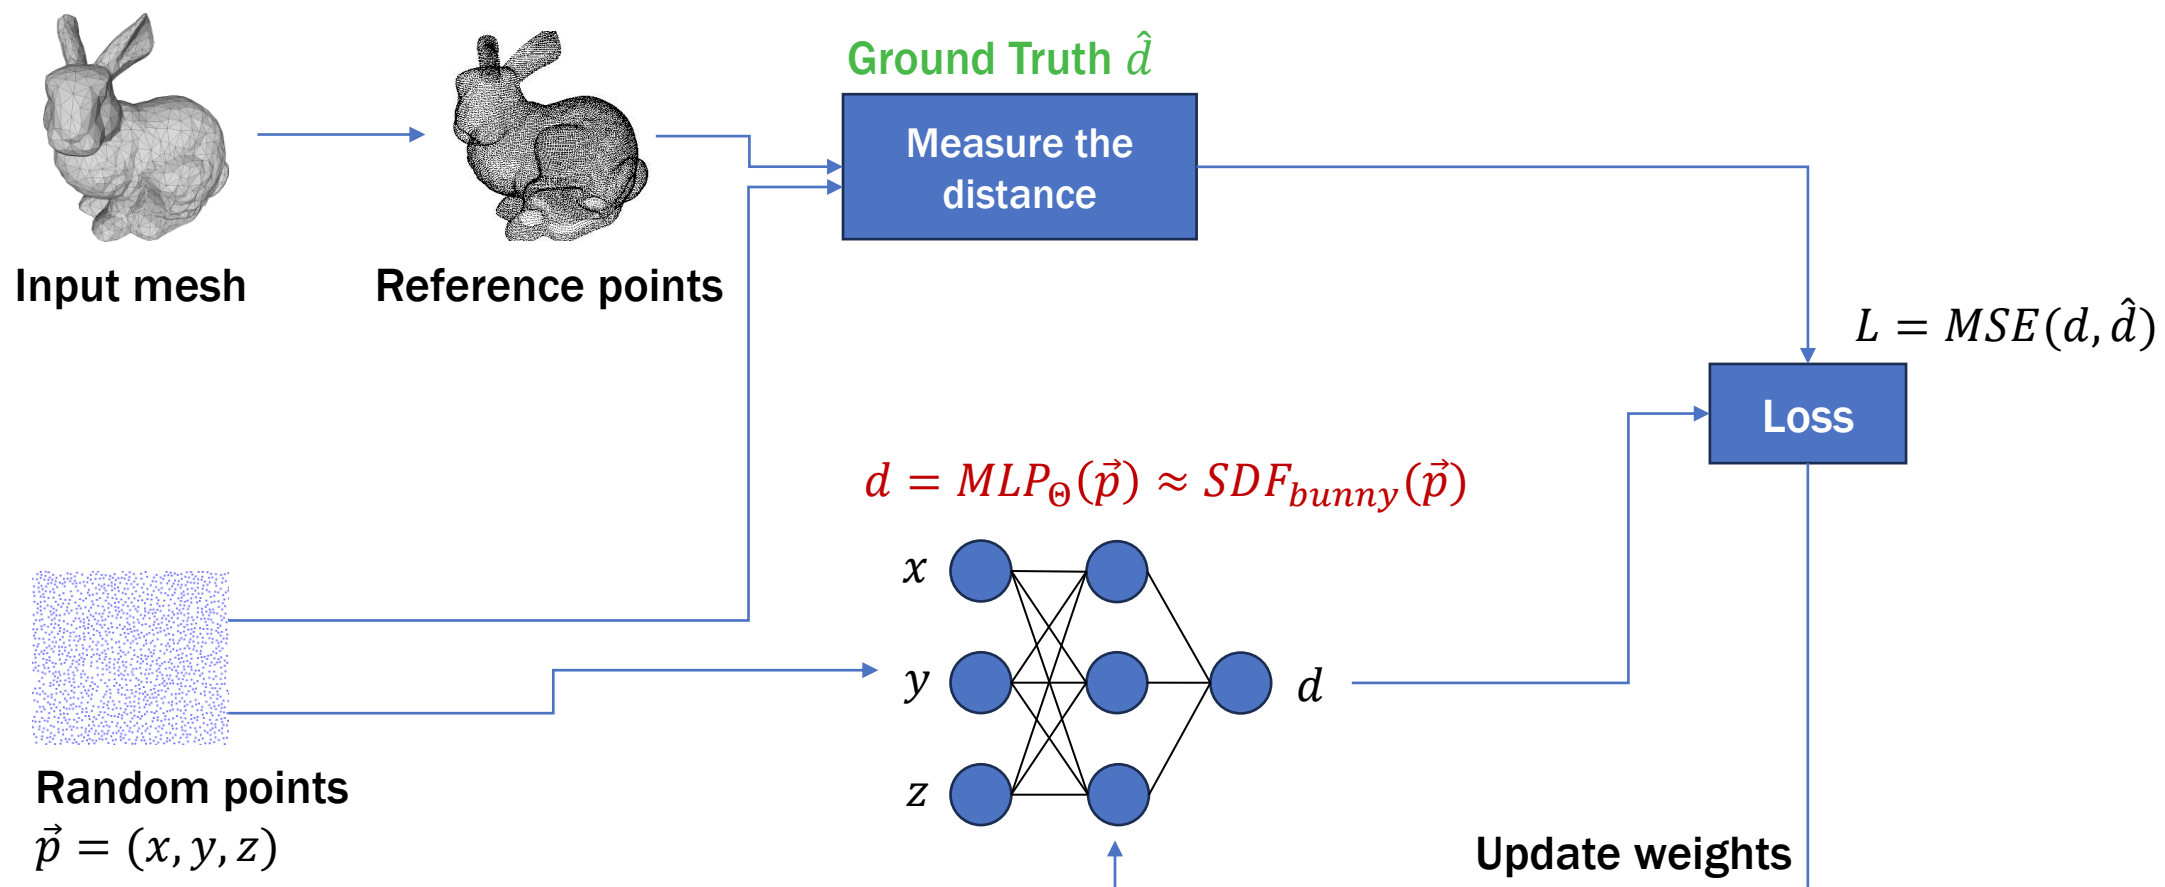
TRAINING A NEURAL SHAPE



TRAINING A NEURAL SHAPE



TRAINING A NEURAL SHAPE



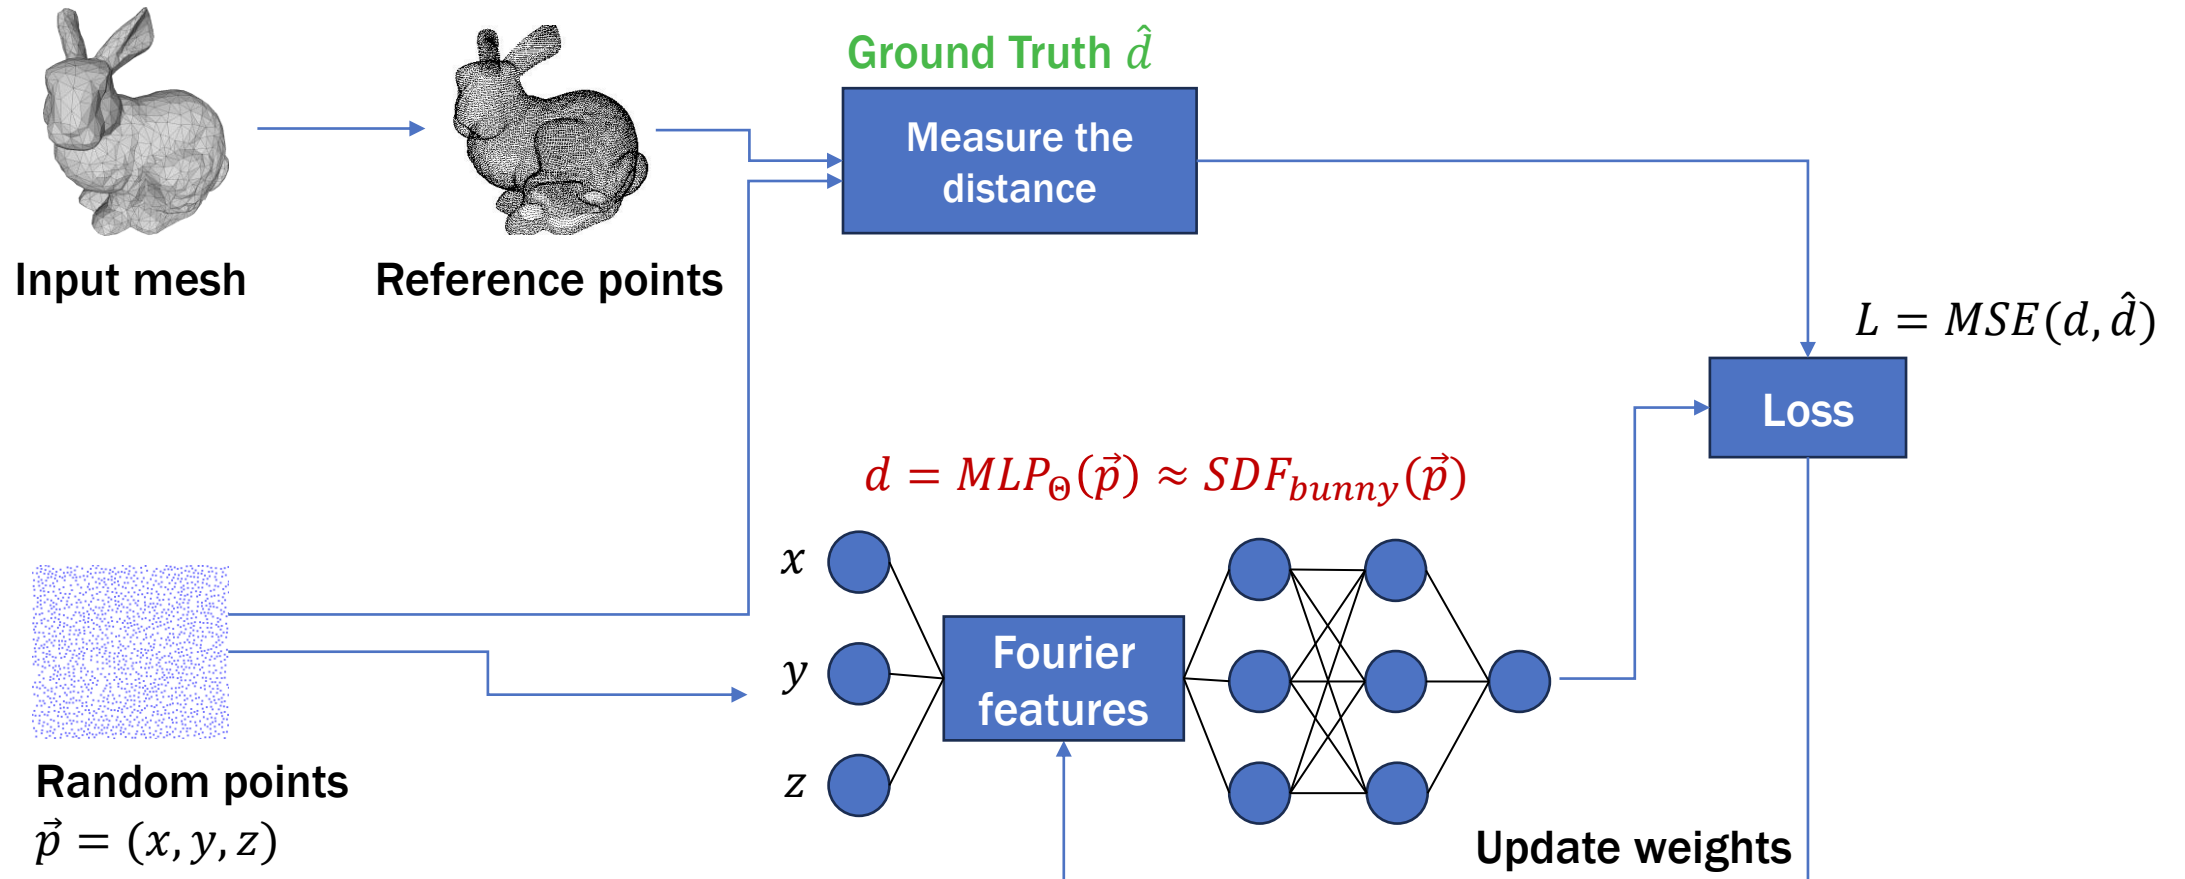
ADDING DETAILS

- We decompose the input vector into a set of frequencies (cosine and sine)
- w_1, \dots, w_n are learned
- The model is forced to extract the high frequencies
 - High frequencies encode fine details (also in images)

$$\phi(\vec{p}) = [\cos(2\pi w_1 \vec{p}), \sin(2\pi w_1 \vec{p}), \dots, \cos(2\pi w_n \vec{p}), \sin(2\pi w_n \vec{p})]^T$$

- Those features are called **Fourier features**
- <https://arxiv.org/abs/2006.10739>

FINAL PIPELINE






MY RESEARCH

EXPLAINABLE GENERATIVE MODELS

Current 3D generative models are **not explainable**



3D artists want to manipulate the generated shapes

- «I want a taller bottle»
- «I want a thicker cup»

APPLYING OPERATORS TO THE SHAPES

Operators

- Implicit representations offer trivial way to combine primitives


$$s_1(p, c, r) = \text{dist}(p, c) - r$$

$$s_1(p, c, r) = \text{dist}(p, c) - r$$

$$s_1 \cup s_2 = \min(s_1, s_2)$$


$$s_1 \cap s_2 = \max(s_1, s_2)$$

$$s_1 \cup s_2 = \min(-s_1, s_2)$$



MORE COMPLEX OPERATIONS

We can blend primitives to get more detail



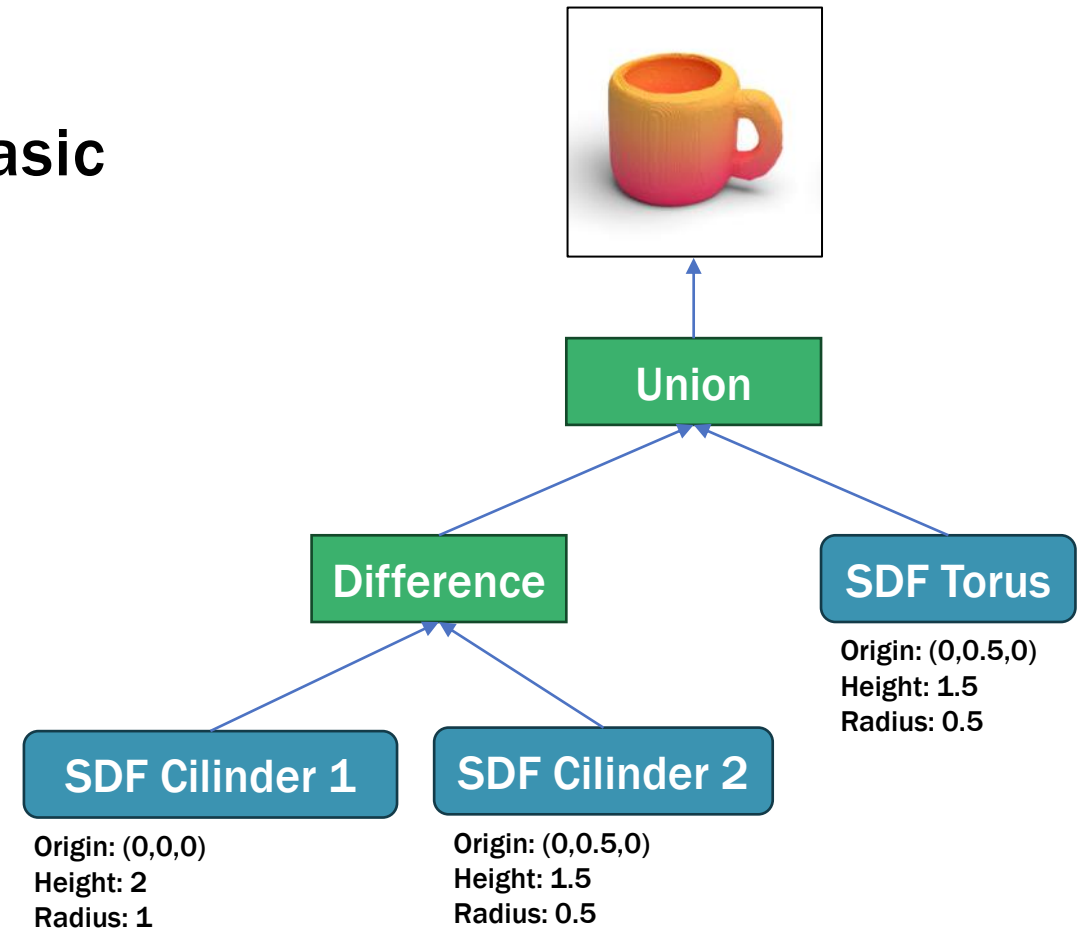

```
float opSmoothUnion( float d1, float d2, float k )
{
    float h = clamp( 0.5 + 0.5*(d2-d1)/k, 0.0, 1.0 );
    return mix( d2, d1, h ) - k*h*(1.0-h);
}

float opSmoothSubtraction( float d1, float d2, float k )
{
    float h = clamp( 0.5 - 0.5*(d2+d1)/k, 0.0, 1.0 );
    return mix( d2, -d1, h ) + k*h*(1.0-h);
}

float opSmoothIntersection( float d1, float d2, float k )
{
    float h = clamp( 0.5 - 0.5*(d2-d1)/k, 0.0, 1.0 );
    return mix( d2, d1, h ) + k*h*(1.0-h);
}
```

CONSTRUCTIVE SOLID GEOMETRY




Create a detailed 3D model using basic building blocks



EXAMPLE TASK: MODEL RECONSTRUCTION



From a given 3D mesh (explicit representation),
predict the constructive solid geometry tree.

The root node is still an SDF, but it's **explainable**.



RESULTS

Progress: 1/2000



An abstract 3D rendering of a torus and a cylinder. The torus is primarily orange and yellow, with a dark, reflective interior. The cylinder is primarily blue and yellow, also with a reflective surface. They are positioned in a way that they appear to be interacting or intersecting. The background is a gradient of dark brown and blue.

NOTEBOOK

https://colab.research.google.com/drive/109t1n8Nv7SiepqEKx8tHJyydF_baxF8H